

EUCALYPTUS

3.0.1 Administration Guide

2012-04-30 Eucalyptus Systems

Contents

Welcome.....	4
Overview of Eucalyptus.....	4
Accessing Eucalyptus.....	4
Command Line Interface.....	4
Eucalyptus Dashboard Overview.....	4
Managing the Cloud.....	9
Cloud Overview.....	9
Best Practices.....	10
Working with vSphere.....	10
Securing Your Cloud.....	11
High Availability.....	12
Cloud Tasks.....	14
Inspect System Health.....	15
Monitor Usage.....	16
View User Resources.....	16
List Arbitrators.....	17
Add a Node.....	17
Remove a Node.....	17
Restart Eucalyptus.....	18
Shut Down Eucalyptus.....	19
Back Up and Restore Eucalyptus.....	21
Managing Access.....	23
Access Overview.....	23
Access Concepts.....	23
Policy Overview.....	25
LDAP/AD Integration.....	34
Access Tasks.....	40
Use Case: Creating an Administrator.....	41
Use Case: Creating a User.....	42
Accounts.....	43
Groups.....	45
Users.....	50
Credentials.....	54
Synchronize LDAP/AD.....	55
Managing Images.....	57
Image Overview.....	58
Image Tasks.....	58
Add an Image.....	58
Modify an Image.....	61
Create an Image.....	61
Associate a Kernel and Ramdisk.....	85
Bundle an Image for Amazon EC2.....	86
Bundle a Windows Instance.....	86
Appendix.....	87
Eucalyptus Administration Commands.....	87
euca_conf.....	87
euca-describe-properties.....	89
euca-describe-services.....	90
Troubleshooting Eucalyptus.....	91
Log Files.....	91

Network Information.....	92
Walrus and Storage.....	92
Access and Identities.....	93
Windows Images.....	93
Instances.....	95
Recovering from a Fail: Walrus.....	95
Glossary.....	98

Welcome

This is the Eucalyptus Administration Guide. The guide shows you how to access Eucalyptus with a web-based GUI and with command line tools. This guide also describes how to perform common administration tasks in three areas: system management, identity management, and resource management.

This document is intended to be a reference. You do not need to read it in order, unless you are following the directions for a particular task.

Overview of Eucalyptus

Eucalyptus is a Linux-based software architecture that implements scalable, efficiency-enhancing private and hybrid clouds within an enterprise's existing IT infrastructure. Because Eucalyptus provides Infrastructure as a Service (IaaS), you can provision your own resources (hardware, storage, and network) through Eucalyptus on an as-needed basis.

A Eucalyptus cloud is deployed across your enterprise's on-premise data center. As a result, your organization has a full control of the cloud infrastructure. You can implement and enforce various level of security. Sensitive data managed by the cloud does not have to leave your enterprise boundaries, keeping data completely protected from external access by your enterprise firewall.

Eucalyptus was designed from the ground up to be easy to install and non-intrusive. The software framework is modular, with industry-standard, language-agnostic communication. Eucalyptus is also unique in that it provides a virtual network overlay that isolates network traffic of different users as well as allows two or more clusters to appear to belong to the same Local Area Network (LAN).

Eucalyptus also is compatible with Amazon's EC2, S3, and IAM services. This offers you hybrid cloud capability.

Accessing Eucalyptus

There are two ways to interact with Eucalyptus. You can use the administrative command line interface for making requests to Eucalyptus, or you can use the web-based user interface, called the Eucalyptus Dashboard.



Tip: This guide will show both CLI and Dashboard steps for performing a task, when the task can be performed by both methods.

Command Line Interface

Eucalyptus supports two command line interfaces (CLIs). The administration CLI is installed when you install Eucalyptus server-side components. The administration CLI is for maintaining and modifying Eucalyptus.

The other user CLI, called Euca2ools, can be downloaded and installed on clients. Euca2ools are for end users and can be used with both Eucalyptus and Amazon Web Services (AWS).

The commands used in this guide assume that the environment variables exported by a eucarc file for an administrative Eucalyptus user have been set. For more information, see the [Eucalyptus Installation Guide](#).

Eucalyptus Dashboard Overview

The Eucalyptus Dashboard provides cloud administrators with a way to perform several management tasks in a web user interface.

The Dashboard provides **Quick Links** for standard administrative actions and queries. These links, located on the left side of the screen, provide a convenient way to navigate through the Dashboard. For example, if you click **Accounts**, the Dashboard displays the **Accounts** page, listing all accounts in your system. Any property of a link, e.g., an account ID, displays on the right side of the screen as a link.

For more experienced users, the Dashboard provides a robust search engine. You can search for information or tasks quickly by building your own search. Because the Dashboard is search-based, even the **Quick Links** and other returned URLs from searches are themselves searches. Because each link is search, any Dashboard link you bookmark is also a search.

So the Dashboard offers you two ways to get information: by search or by following links. To show the member users of an account, you can click **Accounts** in quick links, select the account, and then click on the **Member users** link in the **Properties** section. Or, you use the Search box and type: `user:account=<account_name>`.

Signing in to the Eucalyptus Dashboard

This section describes how to sign in to the Eucalyptus Dashboard.

The Eucalyptus Dashboard is web-based interface that allows you to manage your system, identities, and resources.

To sign in to the Dashboard:

1. Open a browser window and go to `https://<CLC_IP_address>:8443`
Your browser displays a warning.
2. Accept the self-signed SSL certificate and continue.
The Eucalyptus sign-in page displays.
3. Enter your account name in the **Account** field.
 - For system admins, the account name is `eucalyptus`.
4. Enter your user name in the **User** field.
5. Enter your password in the **Password** field.
6. Click the **Sign in** button.
The Dashboard **Start Guide** page displays.

You can now use the Dashboard to manage your system, identities, and resources.

Understanding the Eucalyptus Dashboard

This section explains the components of the Dashboard screen.

The Dashboard screen has the following areas:

Header

This area includes the logo, the link to a user profile setting menu, and the big search box.

User Profile

The current login user identity displays on the right side of the logo area. It shows the user name and the account name of the user identity, in the format of

```
<user_name>@<account_name>
```

Click the profile name to display the user profile menu. The menu provides the following functions:

- **View/change profile:** Displays the search result of the current user. In the search result page, you can view or change your identity's profile.
- **View access key:** Displays the search result of the current user's access key.
- **Change password:** Displays a dialog to change password.

Quick Links

- **Download new credentials:** Downloads the current user's credential package in a zip file.

The left side of the Dashboard screen contains the **Quick Links** area. This area provides links to various contents of the Dashboard.

The **Quick Links** area is organized into sections made up of two levels. The top level is a heading for that section. Under each heading is a second section that contains a list of links. Each link is a search query in the form of the URL. Click a link to return the associated search result. For example, **Your Keys** is a search query of all the access keys belonging to you.

You can hide the **Quick Links** area by clicking the arrow of the vertical separator between **Quick Links** and the main content area.

Main Content

The center part of the main screen displays the main content, usually the search result list. In many content displays, the Dashboard displays a toolbar that contains action buttons. The bottom of the content area provides the page navigation controls.

The search result list usually has multiple columns, some of which are sortable. Click the title in the column to sort the column display. If the list is too long, the Dashboard partitions the list into multiple pages. By default, each page displays a maximum of 25 rows, but you can configure this number.

.

When you select an item in the main content area, the Dashboard highlights the entire row and displays the **Properties** area. To select multiple items, use the **Ctrl** key for individual items, or the **Shift** key for a continuous block of items.

Properties

The **Properties** area displays the detailed information about a selected search result item. The properties are displayed in two columns: the property name is on the left, and the property value is on the right.

Working with the **Properties** area:

- The Dashboard displays values of editable properties in a white input box. If you make any changes to value, the Dashboard displays the **Save** button at the bottom. Click this button to save any changed values.
- Some properties are of complex types. For example, the list of member users of an account. In these cases, the property names are displayed in hyperlinks with a magnifying glass icon. These hyperlinks invoke a search query.
- Other properties display an "action" icon. For example, **Password** displays a pencil icon. Click that icon to change the password.

- The Dashboard allows you to customize the displayed information in **Properties**. Click the plus icon to add a new property to the display. Click the minus icon to delete a property from the display.
- Click the **X** next to the **Properties** title to hide the area.

Status

The bar at the bottom of the main screen shows system status messages, log window toggle button and the software version (from left to right).

Logs

Click the **LOG** button on the status bar to pop up the log window. The log window records important dashboard events, especially any operations that modify system states, e.g. adding a new account, etc. The log windows records the latest 1024 log messages.

Using Search

This section details the **Search** function in the Eucalyptus Dashboard.

Experienced users can use search box to get any information provided by the dashboard. The basic syntax of a search is as follows:

```
<type>: <field1>=<value1>,<value2>
<field2>=<value1>,<value2>
```

The `<type>` specifies the information type provided by the Dashboard. Currently Eucalyptus supports the following types:

Type Name	Description	Fields
start	Start page	None
config	Service components configuration	None
account	Accounts	name , id
group	User groups	account , name , id, path , user
user	Users	account , name , id, path , enabled, registration, group, [custom keys...] <ul style="list-style-type: none"> • User's custom keys can be used as fields. • <code>group</code> field means the user has membership in that group.
policy	Eucalyptus policies	account , user , group , name , id, version , text
key	Access keys	account , user , id, active, user
cert	X509 certificates	account , user , id, revoked, active
image	VM images	None
vmtype	VM types	None
report	Report page	None



Note: In the table, the field names in bold font means that for that field, the query evaluator does a partial match for the value.

The minimal search query contains the type name and a colon. For example, to display the **Start Guide** page, you would enter:

```
start:
```

After the colon, you enter a list of conditions, if any are accepted by the type name. Each condition has a field name and a list of values. The field name and values are separated by an equal sign. There is no space between the field name and value. Separate values with a comma, and don't include a space. Separate multiple conditions with a space.

To evaluate the search query, all conditions must be satisfied. For each condition, only one of the value for the field needs to be matched. For example, to find all users in the accounts whose names contain "testaccount", and whose user names contain "user1" or "user2", and who are enabled, enter the following:

```
user:account=testaccount name=user1,user2 enabled=true
```

After entering a search query in the search box in the header area of the main screen, press the **Enter** key. The search result displays in the content area. The browser URL will also change to reflect the search. Actually, the search query itself is part of the URL (after the pound sign). For example:

```
https://localhost:8443/#account:name=test
```

In fact, you can type a search directly in the URL box of the browser. But remember that the URL itself is URL encoded. This also enables Eucalyptus to construct a search URL and add to any web page.

Managing the Cloud

The **System Management** section of the **Quick Links** area allows you to go to the **Start Guide** or the **Service Components** page.

Cloud Overview

Eucalyptus is comprised of six components: Cloud Controller, Walrus, Cluster Controller, Storage Controller, Node Controller, and an optional VMWare Broker. Each component is a stand-alone web service. This architecture allows Eucalyptus both to expose each web service as a well-defined, language-agnostic API, and to support existing web service standards for secure communication between its components.

Cloud Controller

The Cloud Controller (CLC) is the entry-point into the cloud for administrators, developers, project managers, and end-users. The CLC queries the Node Controller for information about resources, makes high-level scheduling decisions, and makes requests to the Cluster Controllers. As the interface to the management platform, the CLC is responsible for exposing and managing the underlying virtualized resources (servers, network, and storage). You can access the CLC through Amazon's Elastic Compute Cloud (EC2) and through a web-based Dashboard.

Walrus

Walrus allows users to store persistent data, organized as buckets and objects. You can use Walrus to create, delete, and list buckets, or to put, get, and delete objects, or to set access control policies. Walrus is interface compatible with Amazon's Simple Storage Service (S3). It provides a mechanism for storing and accessing virtual machine images and user data. Walrus can be accessed by end-users, whether the user is running a client from outside the cloud or from a virtual machine instance running inside the cloud.

Cluster Controller

The Cluster Controller (CC) generally executes on a machine that has network connectivity to both the machines running the Node Controller (NC) and to the machine running the CLC. CCs gather information about a set of node machines and schedules virtual machine (VM) execution on specific nodes. The CC also manages the virtual machine networks. All nodes associated with a single CC must be in the same subnet.

Storage Controller

The Storage Controller (SC) provides functionality similar to the Amazon Elastic Block Store (Amazon EBS). The SC is capable of interfacing with various storage systems (NFS, iSCSI, SAN devices, etc.). Elastic block storage exports storage volumes that can be attached by a VM and mounted or accessed as a raw block device. EBS volumes persist past VM termination and are commonly used to store persistent data. An EBS volume cannot be shared between VMs and can only be accessed within the same availability zone in which

Node Controller

the VM is running. Users can create snapshots from EBS volumes. Snapshots are stored in Walrus and made available across availability zones. Eucalyptus with SAN support lets you use your enterprise-grade SAN devices to host EBS storage within a Eucalyptus cloud.

The Node Controller (NC) executes on any machine that hosts VM instances. The NC controls VM activities, including the execution, inspection, and termination of VM instances. It also fetches and maintains a local cache of instance images, and it queries and controls the system software (host OS and the hypervisor) in response to queries and control requests from the CC. The NC is also responsible for the management of the virtual network endpoint.

VMware Broker

VMware Broker (Broker) is an optional Eucalyptus component activated only in versions of Eucalyptus with VMware support. Broker enables Eucalyptus to deploy virtual machines (VMs) on VMware infrastructure elements. Broker mediates all interactions between the CC and VMware hypervisors (ESX/ESXi) either directly or through VMware vCenter. For more information about working with vSphere Server, see [Working with vSphere](#).

Best Practices

This section details Eucalyptus best practices for your private cloud.

Working with vSphere

Eucalyptus with the VMware Broker option can create and manage virtual machines on all or a subset of vSphere infrastructure resources. Since Eucalyptus takes over the task of managing Virtual Machines, to avoid interfering with the operation of your cloud, it is important to avoid performing some operations through vSphere-specific tools, such as the vSphere Client. Conversely, since Eucalyptus does not manage vSphere hosts, network, or storage, the administrator will have to continue using vSphere-specific tools for other operations. The following is a comprehensive, but not exhaustive list of operations that should and should not be performed with vSphere-specific tools.

Actions that may be performed with vSphere tools at any time

- vSphere management tasks that do not involve resources (VMs, hosts, networks, datastores, folders, vCenter Server sessions, etc.) that have never been and are not being used by Eucalyptus.
- Adding new resources -- hosts, networks, datastores, folders -- to vCenter. Such resources may be discovered and used by Eucalyptus automatically, either immediately or after a VMware Broker restart, if the VMware Broker configuration allows that (see 'discover' option in the section on configuring the VMware Broker).
- While we recommend using Eucalyptus's API to manage instances, a VM created by Eucalyptus may be powered off using vSphere Server tools. (It is not necessary to delete such a VM as Eucalyptus will delete it.)
- Managing roles and permissions in ways that do not reduce privileges of Eucalyptus since the time it became active.
- Changing vCenter or ESX(i) host settings that do not interfere with ongoing sessions and operations. For instance, a license on a host utilized by Eucalyptus may be changed as long as the host remains operational.
- Changing of vCenter ID or Name (see below regarding the change of IP address).

Actions that may be performed when Eucalyptus is not active (specifically, when the VMware Broker is shut down):

- Deleting the templates (VMs whose names start with 'emi-'). Those will be recreated if needed, albeit at the cost of additional instance start-up delay. To control the space used by and the number of templates, use VMware Broker's configuration properties `vmwarebroker.vsphere_cache_limit_bytes` and `vmwarebroker.vsphere_cache_max_elements`.
- Managing roles and permissions for Eucalyptus-managed resources, as long as new roles, if any, are reflected in VMware Broker configuration (see 'login' parameter) and as long as no Eucalyptus-created running VMs are taken out of Eucalyptus's control.
- vCenter IP address may be changed as long as VMware Broker's configuration is modified accordingly. IP addresses of ESX hosts may be changed as long as there are no running Eucalyptus VMs on the host (furthermore, a change of IP address may require adjustment of configuration unless the host can be discovered).

Actions to avoid with vSphere tools at all times:

- Renaming, modifying the settings for, or cloning Eucalyptus-created inventory objects (the name of the top-level folder on vCenter, VM templates, VMs). This includes changing the virtual hardware characteristics of VMs created by Eucalyptus.
- Migrating, snapshotting, and failing over VMs or templates (emi-...) created by Eucalyptus to a different host or a different datastore with VMotion, VMware HA, or VMware DRS.
- Changing default ports (80 for HTTP and 443 for HTTPS).

Securing Your Cloud

Eucalyptus components receive and exchange messages using either Query or SOAP interfaces (or both). Messages received over these interfaces are required to have some form of a time stamp (as defined by AWS specification) to prevent message replay attacks. Because Eucalyptus enforces strict policies when checking timestamps in the received messages, for the correct functioning of the cloud infrastructure, it is crucial to have clocks constantly synchronized (for example, with `ntpd`) on all machines hosting Eucalyptus components. To prevent user commands failures, it is also important to have clocks synchronized on the client machines.

Following the AWS specification, all Query interface requests containing the Timestamp element are rejected as expired after 15 minutes of the timestamp. Requests containing the Expires element expire at the time specified by the element. SOAP interface requests using WS-Security expire as specified by the WS-Security Timestamp element.

When checking the timestamps for expiration, Eucalyptus allows up to 20 seconds of clock drift between the machines. This is a default setting. You can change this value for the CLC at runtime by setting the `bootstrap.webservices.clock_skew_sec` property as follows:

```
euca-modify-property -p
bootstrap.webservices.clock_skew_sec=<new_value_in_seconds>
```

For additional protection from the message replay attacks, the CLC implements a replay detection algorithm and rejects messages with the same signatures received within 15 minutes.



Important: To protect against replay attacks, the CLC only caches messages for 15 minutes. So it's important that any client tools used to interact with the CLC have the Expires element set to a value less than 15 minutes from the current time. This is usually not an issue with standard tools, such as `euca2ools` and Amazon EC2 API Tools.

You can configure replay detection in the CLC to allow replays of the same message for a set time period. This might be needed to ensure that legitimate requests submitted by automated scripts closely together (such as two requests to describe instances issued within the same second) are not rejected as malicious. The time within which messages with the same signatures are accepted is controlled by the

`bootstrap.webservices.replay_skew_window_sec` property. The default value of this property is three seconds. To change this value, enter the following command:

```
euca-modify-property -p
bootstrap.webservices.replay_skew_window_sec=<new_value_in_seconds>
```

If you set this property to 0, Eucalyptus will not allow any message replays. This setting provides the best protection against message replay attacks, but may break some of the client-side scripts that issue commands too quickly.

If you set this property to any value greater than 15 minutes plus the values of `ws.clock_skew_sec` (that is, to a value ≥ 920 sec in the default installation), Eucalyptus disables replay detection completely.

High Availability

High availability is the result of the combination of functionality provided by Eucalyptus and the environmental and operational support to maintain the constituent systems's proper operation. Eucalyptus provides functionality aimed at enabling highly available operation:

- Detection of service faults and monitoring of system health: gather service status, determine current service topology, admit requests which can be satisfied using only healthy services in that topology
- Tools for interrogating the system's health: access to service state information
- Error gathering to aid in determining the cause: access to fault information as it impacts service function
- Automated failover when redundant services are configured: removal of faulty services and enabling of healthy services
- Service state control: ability to remove individual component-services (when configured with HA pair) from operation without disrupting service
- Replacement/restoration of component-services: procedures for restoring/replacing a component service after a total-loss failure (e.g., disk failure, host combustion, etc.)

In addition to previously detailed deployment recommendations, delivering highly available services with Eucalyptus depends on appropriate operational and maintenance support. The following sections detail the related system functionality and procedures.

Understanding Service State

The system monitors service health and enables healthy services to process user requests while marking faulty services as being NOTREADY. Each component service is interrogated by the system to determine its current state. Faults are detected either:

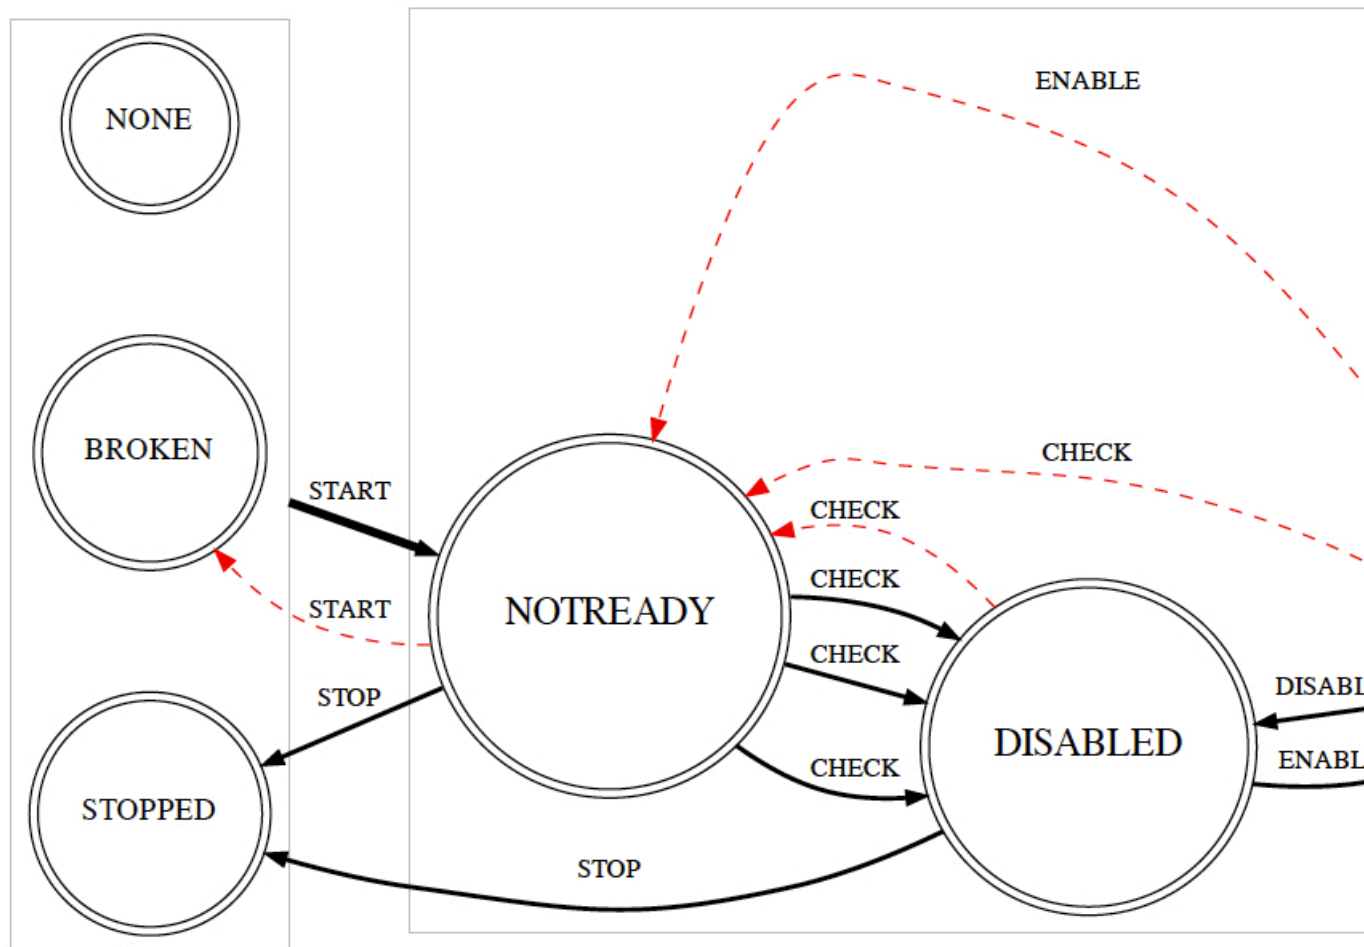
- The service reporting a fault has been detected (for example, due to misconfiguration, dependency service failure, environmental fault, etc.)
- Failure to contact the service

The following table overviews the relevant states

State	Operational	In use by system	Description
ENABLED	Yes	Yes	Service is operating correctly and is selected for processing requests
DISABLED	Yes	No	Service is operating correctly but is not selected for processing requests
NOTREADY	No	No	Service is failing to operate correctly
BROKEN	No	No	Service is not contactable by the system

State	Operational	In use by system	Description
STOPPED	N/A	No	Service has been stopped by an administrator

The following diagram indicates the set of relevant states and transitions between them. Black arrows indicate a transition between states that is initiated by the system or an administrator request. Red errors indicate a failure to transition into the originating state that results in a transition to the destination error state.



Based on the collected service state, the system will:

- Attempt to advance previously non-functioning services to a functional state
- Determine whether any functioning services can be **ENABLED** and added to the set used for serving requests

Understanding System Availability

The impact of a service fault on the system's availability depends upon the deployment and configuration of the system. The following table details the scope a service fault can have on system availability for each component type.

Component	Scope (Fault Region)	Description
Cloud Controller	Cloud	The CLC is a cloud-wide service and must have at least one operation service.

Component	Scope (Fault Region)	Description
Walrus	Cloud	Walrus is a cloud-wide service and must have at least one operation service.
Cluster Controller	Availability Zone	CCs are associated with a partition and service requests specific to an availability zone. Should an availability zone not have an operational CC, instance requests will be rejected for the corresponding zone.
Storage Controller	Availability Zone	Storage controllers are associated with a partition and service requests specific to an availability zone. Should an availability zone not have an operational storage controller volume and snapshot creation requests will be rejected for the corresponding zone
VMware Broker	Availability Zone	VMware Broker are associated with a partition and service requests specific to an availability zone. Should an availability zone not have an operational VMware Broker instance requests will be rejected for the corresponding zone
Arbitrators	User-facing Service Host	Arbitrators are associated with a host that runs user-facing component services (CLC, Walrus). Each host must have an operational Arbitrator. Should a component service host have a configured but faulty Arbitrator, a fail-stop condition occurs and locally hosted services report a NOTREADY error.
Node Controller	Compute Host	NCs are associated with each node and interact with the hypervisor to service node-specific requests.

A quick way to evaluate system availability is to determine whether:

- The cloud has an enabled CLC
- The cloud has an enabled Walrus
- The availability zone has an enabled CC
- The availability zone has an enabled SC
- The user-facing service host has one reachable Arbitrator per host (if you configure an Arbitrator)

Cloud Tasks

Listing of the cloud-related tasks.

You can perform the following cloud-related tasks listed in this section:

- [Inspect System Health](#)
- [Monitor Usage](#)
- [List Arbitrators](#)
- [Add a User](#)
- [Delete a User](#)
- [Add a Node](#)
- [Remove a Node](#)
- [Restart Eucalyptus](#)
- [Shut Down Eucalyptus](#)
- [Back Up and Restore Eucalyptus](#)

Inspect System Health

Eucalyptus provides access to the current view of service state and the ability to manipulate the state. You can inspect the service state to either ensure system health or to identify faulty services. You can modify a service state to maintain activities and apply external service placement policies.

View Service State

Use the `euca-describe-services` command to view the service state. The output indicates:

- Component type of the service
- Partition in which the service is registered
- Unique name of the service
- Current view of service state
- Last reported epoch (this can be safely ignored)
- Service URI
- Fully qualified name of the service (This is needed for manipulating services that did not get unique names during registration. For example: internal services like reporting or DNS)

The default output includes the services that are registered during configuration, as well as information about the DNS service, if present. You can obtain additional service state information, such as internal services, by providing the `-system-internal` flag.

You can also make requests to retrieve service information that is filtered by either:

- current state (for example, NOTREADY)
- host where service is registered
- partition where service is registered
- type of service (for example, CC or Walrus)

When you investigate service failures, you can specify `-events` to return a summary of the last fault. You can retrieve extended information (primarily useful for debugging) by specifying `-events -events-verbose`.

Modify Service State

To modify a service:

Enter the following command:

```
euca-modify-service <service_name> -s <new_service_state>
```

If, for example, you have SCs that are correctly configured and operating in HA mode. However, you want to shut down the primary SC for maintenance. The primary SC is SC00 and the secondary SC is SC01. SC00 is `ENABLED` and SC01 is `DISABLED`.

To disable SC00 and cause SC01 to take over, you would enter:

```
euca-modify-service SC00 -s DISABLED
```

To check status of services, you would enter:

```
euca-describe-services
```

When SC01 is marked **ENABLED**, the eucalyptus-cloud process on the host that SC00 is running can be safely shutdown and maintenance tasks can be performed. When maintenance is complete, you can start the eucalyptus-cloud process on SC00. SC00 will enter the **DISABLED** state by default. You can chose to let SC00 continue to be the primary and SC01 will be the secondary.

If you want to designate SC00 as the primary, make sure no volumes or snapshots are being created and that no volumes are being attached or detached, and then enter:

```
euca-modify-service SC01 -s DISABLED
```

Monitor the state of services using `euca-describe-services` until SC01 is marked **DISABLED** and SC00 is **ENABLED**.

Monitor Usage

Generate reports to monitor cloud resource use over a specified date range. Usage can be reported for instances, storage, and Walrus.

- **Instances:** Shows a count and duration of instances running, sorted by image size, as well as I/O net and disk, during the date range specified. If desired, the report can be displayed on a per-user, per-account, per-cluster or per-availability zone basis. Once filtered, reports can be grouped using the same criteria as filters. For example, you might display the results by user, and then group the results so that they are displayed based on accounts, e.g. each user in account A is shown followed by each user in account B, etc.
- **Storage:** Shows GB of storage used (max and GB-days) for volumes and snapshots during the date range specified. Display filters and grouping are the same as for instances.
- **Walrus:** Shows maximum buckets and objects used through Walrus during the date range specified. Display filters and grouping are available on a per-user or per-account basis.

You can generate reports in either CSV, PDF, or HTML formats for use with external tools.

To generate reports:

1. Log in to the Dashboard at https://<CLC_hostname_or_ipaddress>:8443.
2. Click **Usage Report** in the **Quick Links** section.
The **Usage Report** page displays.
3. In the **Start date** field, select the date to start
4. Click the **Generate** button.

The usage report displays, reflecting your choices in the report.

View User Resources

To see resource use by your cloud users, Eucalyptus provides the following commands with the `-verbose` flag:

- `euca-describe-groups verbose`: Returns information about security groups in your account, including output type identifier, security group ID, security group name, security group description, output type identifier, account ID of the group owner, name of group granting permission, type of rule, protocol to allow, start of port range, end of port range, source (for ingress rules) or destination (for egress rules), and any tags assigned to the security group.

- `euca-describe-instances verbose`: Returns information about your instances, including output type identifier, reservation ID, name of each security group the instance is in, output type identifier, instance ID for each running instance, EMI ID of the image on which the instance is based, public DNS name associated with the instance (for instances in the running state), private DNS name associated with the instance (for instances in running state), instance state, key name, launch index, instance type, launch time, availability zone, kernel ID, ramdisk ID, monitoring state, public IP address, private IP address, type of root device (ebs or instance-store), placement group the cluster instance is in, virtualization type (paravirtual or hvm), any tags assigned to the instance, hypervisor type, block device identifier for each EBS volume the instance is using, along with the device name, the volume ID, and the timestamp.
- `euca-describe-keypairs verbose`: Returns information about key pairs available to you, including keypair identifier, keypair name, and private key fingerprint.
- `euca-describe-snapshots verbose`: Returns information about EBS snapshots available to you, including snapshot identifier, ID of the snapshot, ID of the volume, snapshot state (pending, completed, error), timestamp when snapshot initiated, percentage of completion, ID of the owner, volume sized, description, and any tags assigned to the snapshot.
- `euca-describe-volumes verbose`: Describes your EBS volumes, including volume identifier, volume ID, size of the volume in GiBs, snapshot from which the volume was created, availability zone, volume state (creating, available, in-use, deleting, deleted, error), timestamp of the volume creation, and any tags assigned to the volume.

List Arbitrators

To see a list of Arbitrators running on your cloud:

- Enter the following command to display Arbitrators for the current CLC or Walrus:

```
/usr/sbin/euca-describe-services --system-internal
```

- Enter the following command to display Arbitrators on both primary and secondary CLCs or Walruses:

```
/usr/sbin/euca_conf --list-arbitrators
```

Add a Node

Describes how to add a node to your system.

If you want to increase your system's capacity, you'll want to add more nodes (NC servers). To add a node, perform the following tasks:

1. Log in to the CLC and enter the following command:

```
/usr/sbin/euca_conf --register-nodes \ "<nodeName1> ...  
<nodeNameN>"
```

2. When prompted, enter the password to log into <nodeNameN>.

Eucalyptus requires this password to propagate the cryptographic keys.

Remove a Node

Describes how to delete nodes in your system.

If you want to decrease your system's capacity, you'll need to decrease nodes (NC servers). To delete a node, perform the following tasks.

Log in to the CLC and enter the following command:

```
/usr/sbin/euca_conf --deregister-nodes "<nodeName1> ... <nodeNameN>"
```

Restart Eucalyptus

Describes the recommended processes to restart Eucalyptus, including terminating instances and restarting Eucalyptus components.

You must restart Eucalyptus whenever you make a physical change (e.g., switch out routers), or edit the `eucalyptus.conf` file. To restart Eucalyptus, perform the following tasks in the order presented.



Note: Before you restart Eucalyptus, we recommend that you notify all users that you are terminating all instances.

1. *Shut Down All Instances*
2. *Restart the CLC*
3. *Restart Walrus*
4. *Restart the CC*
5. *Restart the SC*
6. *Restart a NC*

Shut Down All Instances

To terminate all instances on all nodes:

Enter the following command:

```
euca-terminate-instances <instance_id>
```

Restart the CLC

Log in to the CLC and enter the following command:

```
service eucalyptus-cloud restart
```

All Eucalyptus components on this server will restart.

Restart Walrus

Log in to Walrus and enter the following command:

```
service eucalyptus-cloud restart
```

Restart the CC

You can start the CC two different ways. A normal restart causes the CC to re-parse the `eucalyptus.conf` file without resetting the networking state of the cloud. This means that the CC maintains routes, firewall rules, IP assignments for running VMs. This is how you restart under normal circumstances.

A clean restart causes the CC to reset the networking configuration it maintains. The CC clears any existing network configuration regardless of whether it is in use. That is, a clean restart of a CC managing active VMs can cause the network connectivity they are using to be irreparably lost. We recommend that you only perform this type of restart when you change network modes.

Restart the CC (Normal)

To perform a normal restart:

Log in to the CC and enter the following command:

```
service eucalyptus-cc restart
```

Restart the CC (Clean)

To perform a clean restart:

Log in to the CC and enter the following command:

```
service eucalyptus-cc cleanrestart
```

Restart the SC

Log in to the SC and enter the following command:

```
service eucalyptus-cloud restart
```

Restart a NC

1. Log in to the NC and enter the following command:

```
service eucalyptus-nc restart
```

2. Repeat for each NC.

You can automate the restart command for all of your NCs. Store a list of your NCs in a file called `nc-hosts` that looks like:

```
nc-host-00
nc-host-01
...
nc-host-nn
```

To restart all of your NCs, run the following command:

```
cat nc-hosts | xargs -i ssh root@{ } service eucalyptus-nc
restart
```

Shut Down Eucalyptus

Describes the recommended processes to shut down Eucalyptus.

There may be times when you need to shut down Eucalyptus. This might be because of a physical failure, topological change, backing up, or making an upgrade. We recommend that you shut down Eucalyptus components in the reverse order of how you started them. To stop the system, shut down the components in the order listed.



Note: Before you shut Eucalyptus down, we recommend that you notify all users that you are terminating all instances.

1. *Shut Down All Instances*
2. *Shut Down the NCs*
3. *Shut Down the CCs*
4. *Shut Down the Broker*
5. *Shut Down the SCs*
6. *Shut Down Walrus*
7. *Shut Down the CLC*

Shut Down All Instances

To terminate all instances on all nodes:

Enter the following command:

```
euca-terminate-instances <instance_id>
```

Shut Down the NCs

To shut down the NCs:

1. Log in as root to a machine hosting an NC.
2. Enter the following command:

```
service eucalyptus-nc stop
```

3. Repeat for each machine hosting an NC.

Shut Down the CCs

To shut down the CCs:

1. Log in as root to a machine hosting a CC.
2. Enter the following command:

```
service eucalyptus-cc cleanstop
```



Note: The cleanstop command causes the CC to reset and/or flush all of the existing networking state.

3. Repeat for each machine hosting a CC.

Shut Down the Broker

If your system uses the optional Broker component of Eucalyptus, shut it down by performing the following steps:

1. Log in as root to the machine running both the CC and the VMware Broker.
2. Enter the following command:

```
service eucalyptus-cloud stop
```



Tip: The eucalyptus-cloud stop command also shuts down a CLC, Walrus, and SC components co-located with the CC and VMware Broker to stop at the same time, in the correct order.

Shut Down the SCs

To shut down the SC:

1. Log in as root to the physical machine that hosts the SC.
2. Enter the following command:

```
service eucalyptus-cloud stop
```

3. Repeat for any other machine hosting an SC.

Shut Down Walrus

To shut down Walrus:

1. Log in as root to the physical machine that hosts Walrus.
2. Enter the following command:

```
service eucalyptus-cloud stop
```

Shut Down the CLC

To shut down the CLC:

1. Log in as root to the physical machine that hosts the CLC.
2. Enter the following command:

```
service eucalyptus-cloud stop
```

Back Up and Restore Eucalyptus

You can back up and restore Eucalyptus by completing the following tasks:

1. [Shut Down Eucalyptus](#)
2. [Back Up Eucalyptus](#)
3. [Restore Eucalyptus from a Backup](#)

Shut Down Eucalyptus

Describes the recommended processes to shut down Eucalyptus.

There may be times when you need to shut down Eucalyptus. This might be because of a physical failure, topological change, backing up, or making an upgrade. We recommend that you shut down Eucalyptus components in the reverse order of how you started them. To stop the system, shut down the components in the order listed.



Note: Before you shut Eucalyptus down, we recommend that you notify all users that you are terminating all instances.

1. [Shut Down All Instances](#)
2. [Shut Down the NCs](#)
3. [Shut Down the CCs](#)
4. [Shut Down the Broker](#)
5. [Shut Down the SCs](#)
6. [Shut Down Walrus](#)
7. [Shut Down the CLC](#)

Back Up Eucalyptus

Backing up Eucalyptus involves saving contents of the specific directories on each component. On the CLC host, these directories are:

- The configuration file (/etc/eucalyptus.conf)
- The database files (/var/lib/eucalyptus/db)
- The cryptographic keys (/var/lib/eucalyptus/keys)

On the Walrus host (which, depending on your set-up, might be the same as the CLC host):

- The Walrus buckets (the default location is /var/lib/eucalyptus/bukkits). You should only do this if you have enough spare disk space available. If the Walrus buckets location is an external storage location (like NFS), consult with your storage administrator to back up that location.

On the Storage Controller host (may be the same as the CLC host):

- the SC volumes (the default location is `/var/lib/eucalyptus/volumes`). You should only back this up if you are not using SAN support. If you are using a supported SAN device, consult with your storage administrator to back up volumes stored on the SAN.

To back up Eucalyptus:

1. Calculate the disk space required to store the files about to be backed up (this is most relevant for buckets and volumes, which can be large). For example, on a single-cluster installation with default buckets and volumes paths:

```
du -sh /var/lib/eucalyptus/
```

2. Create a directory for storing these (\$BACKUP) in a location with enough disk space

Restore Eucalyptus from a Backup

In case your system experiences a failure, you can restore Eucalyptus from a backup by performing the following tasks:

1. Shut down any running Eucalyptus components. For further information about the shut down process, see [Shut Down Eucalyptus](#).
2. If you are trying to recover from a broken upgrade by rolling back or by trying the upgrade again, this would be the right time to:
 - Remove all software components related to Eucalyptus (for example, `rpm -e` or `apt-get remove`).
 - Install the appropriate version. For more information, go to the Eucalyptus Installation Guide.



Warning: DEBs will restart the services. Be sure you stop them again before copying back the backed-up files.

3. Replace the saved state.

- If you backed up with `cp`, enter the following:

```
cp -a $BACKUP/eucalyptus.conf /etc/eucalyptus
cp -a $BACKUP/keys $BACKUP/db $BACKUP/bukkits $BACKUP/volumes /var/lib/
eucalyptus
```

- If you backed up with `tar`, enter the following:

```
tar xvf $BACKUP/eucalyptus-backup.tar
```

Managing Access

Eucalyptus manages access control through an authentication, authorization, and accounting system. This system manages user identities, enforces access controls over resources, and provides reporting on resource usage as a basis for auditing and managing cloud activities. The user identity organizational model and the scheme of authorizations used to access resources are based on and compatible with the AWS Identity and Access Management (IAM) system, with some Eucalyptus extensions provided that support ease-of-use in a private cloud environment.

You can also perform user authentication by integrating Eucalyptus with an existing LDAP or Active Directory. In this case, the user, group and account information, and Dashboard login authenticate using the LDAP/AD service. This information cannot be changed from Eucalyptus side when LDAP/AD integration is turned on. However, other Eucalyptus-specific information about user, group and account is still stored within the local database of Eucalyptus, including certificates, secret keys and attached policies.

For more information about synchronizing an existing LDAP or Active Directory with Eucalyptus, see [LDAP/AD Integration](#).

Access Overview

The Eucalyptus design of user identity and access management provides layers in the hierarchical organization of user identities. This gives you refined control over resource access. The core of this design is compatible with the AWS Identity and Access Management (IAM) service. There are also a few Eucalyptus-specific extensions that meet the needs of enterprise customers.

Differences in Access from Eucalyptus 2

The concept of “user” has changed from when Amazon first introduced IAM. The original “user” has become the “account,” and the new “user” is an identity within the “account”. We made a similar change in the current Eucalyptus implementation of an IAM-compatible identity management system. If you upgrade of Eucalyptus 2 to Eucalyptus 3 is performed, the “users” in Eucalyptus 2 are converted into “accounts” in Eucalyptus 3. Please refer to the Eucalyptus upgrade section in the [Installation Guide](#) for the details about this identity conversion process.

In Eucalyptus 2, the access management ability was limited:

- The admin user had full access control of the system resources.
- Regular users controlled resources that they created.
- Certain resources could be shared among users, using mechanisms like image launch permissions and Walrus bucket ACLs.

Eucalyptus 3 introduces a two-tier hierarchy of user identities based on accounts. The access control, therefore, is provided at both tiers:

- Eucalyptus 2 style resource sharing is still available, but now implemented at account level.
- Within each account, fine-grained access control is provided by a policy that is fully compatible with IAM.

Access Concepts

This section describes what Eucalyptus access is and what you need to know about it so that you can configure access to your cloud.

User Identities

In Eucalyptus, user identities are organized into accounts. An account is the unit of resource usage accounting, and also a separate namespace for many resources (security groups, key pairs, users, etc.)

Accounts are identified by either a unique ID (UUID) or a unique name. The account name is equivalent to IAM's account alias. It is used to manipulate accounts in most cases. However, to be compatible with AWS, the EC2 commands often use account ID to display resource ownership. There are command line tools to discover the correspondence of account ID and account name. For example, `euare-accountlist` lists all the accounts with both their IDs and names.

An account can have multiple users, but a user can only be in one account. Within an account, users can be associated with Groups. Group is used to attach access permissions to multiple users. A user can be associated with multiple groups. Because an account is a separate name space, user names and group names have to be unique only within an account. Therefore, user X in account A and user X in account B are two different identities.

Both users and groups are identified by their names, which are unique within an account (they also have UUIDs, but rarely used).

Special Identities

Eucalyptus has two special identities for the convenience of administration and use of the system.

- The **eucalyptus** account: This is like the root group in a Linux system. Every user in the eucalyptus account has full access to the resources in the system. These users are often referred to as sys admin or cloud admin. This account is automatically created when the system starts for the first time. You can not remove eucalyptus account from the system.
- The **admin** user of an account: Each account, including the eucalyptus account, has a user named admin by default, i.e. created automatically by the system when an account is created. The admin of an account has full access to the resources owned by the account. You can not remove the admin user from an account. The admin can delegate resource access to other users in the account by using policies.

Credentials

Each user has a unique set of credentials. These credentials are used to authenticate access to resources. There are three types of credentials:

- An **X.509 certificate** is used to authenticate requests to the SOAP API service.
- A **secret access key** is used to authenticate requests to the REST API service.
- A **login password** is used to authenticate the Eucalyptus Dashboard access.

You can manage credentials using the command line tools (the `euare-` commands) or the Dashboard.

In IAM, each account has its own credentials. In Eucalyptus, the equivalent of account credentials are the credentials of admin user of the account.

You can download the full set of credentials for a user or an account, including X509 certificate and secret access key, by:

```
/usr/sbin/euca_conf --get-credentials
```

or:

```
euca-get-credentials
```

or by using the **Download new credentials** in the Eucalyptus Dashboard.

Whichever way to request the credentials, Eucalyptus returns the following:

- An arbitrary existing active secret access key
- A newly generated X509 certificate

Account Creation

You can create accounts two ways:

- Direct creation using command line tool or Eucalyptus dashboard by sys admin. The accounts created in this method will be available for accessing immediately.
- Registration process. One can apply for an account through the dashboard. The process is as follows:
 1. The cloud user registers using the form on the dashboard interface.
 2. An email will then be sent to the sys admin for review. Sys admin can approve or reject the application by invoking a URL included in the email. A notification email will be sent to the intended account owner.
 3. If the application is approved, the account owner needs to invoke the URL included in the notification email to confirm the approval.
 4. Once confirmed, the new account becomes available.

The account registration status can be found in Dashboard. The account registration status is actually associated with the account admin user. That means you can use the following command to obtain the same information:

```
euare-usergetattributes --delegate account -u admin --show-extra
```

Where the `--show-extra` option displays extra information of a user in the following order:

- Enabled status
- Registration status
- Password expires

The account registration status has the following values based on the status of registration process: REGISTERED, APPROVED, or CONFIRMED. An account that is not confirmed cannot be used or accessed. The system administrator can manipulate the account registration status in both dashboard and command line:

```
euare-usermod --delegate account -u admin --reg-status=status
```

The command line manipulation of the registration status does not send the notification emails. So unless you experienced, we recommend that you use the Dashboard.

Special User Attributes

Eucalyptus extends the IAM model by providing the following extra attributes for a user.

- **Registration status:** This is only meaningful for the account administrator (that is, the account level).
- **Enabled status:** . Use this attribute to temporarily disable a user.
- **Password expiration date**
- **Custom information:** Add any name-value pair to a user's custom information attribute. This is useful for attaching pure text information, like an address, phone number, or department. This is especially helpful with external LDAP or Active Directory services.

You can retrieve and modify the registration status, enabled status, and password expiration date using the `euare-usergetattributes` and `euare-usermod` commands. You can retrieve and modify custom information using `euare-usergetinfo` and `euare-userupdateinfo` commands. For more information, see the [CLI Reference Guide](#) for details about these commands.

Policy Overview

Eucalyptus uses the same policy language to specify user level permissions as AWS IAM. Policies are written in JSON. Each policy file can contain multiple statements, each specifying a permission. A permission statement specifies whether to allow or deny a list of actions to be performed on a list of resources, under specific conditions.

A permission statement has the following components:

- **Effect:** Begins the decision that applies to all following components. Either: "Allow" or "Deny"
- **Action or NotAction:** Indicates service-specific and case-sensitive commands. For example: "ec2:RunInstances"

- **Resource or NotResource:** Indicates selected resources, each specified as an Amazon resource name (ARN). For example: "arn:aws:s3:::acme_bucket/blob"
- **Condition:** Indicates additional constraints of the permission. For example: "DateGreaterThan"

The following policy example contains a statement that gives a user with full permission. This is the same access as the account administrator:

```
{
  "Version": "2011-04-01",
  "Statement": [{
    "Sid": "1",
    "Effect": "Allow",
    "Action": "*",
    "Resource": "*"
  }]
}
```

For more information about policy language, go to the [IAM User Guide](#).

Policy Notes

You can combine IAM policies with account level permissions. For example, the admin of account A can give users in account B permission to launch one of account A's images by changing the image attributes. Then the admin of account B can use IAM policy to designate the users who can actually use the shared images.

You can attach IAM policies to both users and groups. When attached to groups, a policy is equivalent to attaching the same policy to the users within that group. Therefore, a user might have multiple policies attached, both policies attached to the user, and policies attached to the group that the user belongs to.

Do not attach IAM policies (except quota policies, a Eucalyptus extension) to account admins. At this point, doing so won't result in a failure but may have unexpected consequences.

Policy Extensions

Eucalyptus extends the IAM policy in order to meet the needs of enterprise customers.

EC2 Resource

In IAM, you cannot specify EC2 resources in a policy statement except a wildcard, "*". So, it is not possible to constrain a permission on specific EC2 entities. For example, you can't restrict a user to run instances on a specific image or VM type. To solve that, Eucalyptus created the EC2 resource for the policy language. The following example shows the ARN of an EC2 resource.

```
arn:aws:ec2::<account_id>:<resource_type>/<resource_id>
```

Where account id is optional.

Eucalyptus supports the following resource types for EC2:

- image
- securitygroup
- address (either an address or address range: 192.168.7.1-192.168.7.255)
- availabilityzone
- instance
- keypair
- volume
- snapshot
- vmtype

The following example specifies permission to launch instances with only an m1.small VM type:

```
{
  "Version": "2011-04-01",
  "Statement": [{
    "Sid": "2",
    "Effect": "Allow",
    "Action": "ec2:RunInstances",

    "Resource": [
      "arn:aws:ec2::vmtypem1.small",
      "arn:aws:ec2::image/*",
      "arn:aws:ec2::securitygroup/*",
      "arn:aws:ec2::keypair/*",
      "arn:aws:ec2::availabilityzone/*",
      "arn:aws:ec2::instance/*",
    ]
  }]
}
```

Policy Key

Eucalyptus implements the following AWS policy keys:

- aws:CurrentTime
- aws:SourceIp

Eucalyptus extends the policy keys by adding the following to the lifetime of an instance:

- ec2:KeepAlive: specifies the length of time that an instance can run
- ec2:ExpirationTime: specifies the expiration time for an instance

The following example restricts an instance running time to 24 hours:

```
{
  "Version": "2011-04-01",
  "Statement": [{
    "Sid": "3",
    "Effect": "Allow",
    "Action": "ec2:RunInstances",
    "Resource": "*",
    "Condition": {
      "NumericEquals": {
        "ec2:KeepAlive": "1440"
      }
    }
  }]
}
```

If there are multiple `ec2:KeepAlive` or `ec2:ExpirationTime` keys that match a request, Eucalyptus chooses the longer lifetime for the instance to run.

Default Permissions

Different identities have different default access permissions. When no policy is associated with them, these identities have the permission listed in the following table.

Identity	Permission
System admin	Access to all resources in the system

Identity	Permission
Account admin	Access to all account resources, including those shared resources from other accounts like public images and shared snapshots
Regular user	No access to any resource

For convenience, Eucalyptus grants the following default access to regular users:

- Users can list themselves (euare-userlistbypath)
- Users can get their own attributes (euare-usergetattributes)
- Users can get information about themselves (euare-usergetinfo)
- Users can list their own accounts (euare-accountlist)

Account administrators have the following default permissions:

- euare-accountlistpolicies
- euare-accountgetpolicy

Quotas

Eucalyptus adds quota enforcement to resource usage. To avoid introducing another configuration language into Eucalyptus, and simplify the management, we extend the IAM policy language to support quotas. The only addition added to the language is the new `limit` effect. If a policy statement's effect is `limit`, it is a quota statement.

A quota statement also has action and resource fields. You can use these fields to match specific requests, for example, quota only being checked on matched requests. The actual quota type and value are specified using special quota keys, and listed in the `condition` part of the statement. Only condition type `NumericLessThanEquals` can be used with quota keys.



Important: An account can only have a quota policy. If you attach an IAM policy to an account (where, for example, the Effect is "Allow" or "Deny"), there will be unexpected results.

The following quota policy statement limits the attached user to only launch a maximum of 16 instances in an account.

```
{
  "Version": "2011-04-01",
  "Statement": [ {
    "Sid": "4",
    "Effect": "Limit",
    "Action": "ec2:RunInstances",
    "Resource": "*",
    "Condition": {
      "NumericLessThanEquals": {
        "ec2:quota-vminstancenumber": "16"
      }
    }
  } ]
}
```

You can attach quotas to both users and accounts, although some of the quotas only apply to accounts. Quota attached to groups will take no effect.

When a quota policy is attached to an account, it actually is attached to the account administrator user. Since only system administrator can specify account quotas, the account administrator can only inspect quotas but can't change the quotas attached to herself.

The following is all the quota keys implemented in Eucalyptus:

Quota Key	Description	Applies to
s3:quota-bucketnumber	Number of S3 buckets	account and user
s3:quota-bucketobjectnumber	Number of objects in each bucket,	account and user
s3:quota-bucketsize	Size of bucket, in MB	account and user
s3:quota-buckettotalsize	total size of all buckets, in MB	account and user
ec2:quota-addressnumber	Number of elastic IPs	account and user
ec2:quota-imagenummer	Number of EC2 images	account and user
ec2:quota-snapshotnumber	Number of EC2 snapshots	account and user
ec2:quota-vminstancenummer	Number of EC2 instances	account and user
ec2:quota-volumenumber	Number of EC2 volumes	account and user
ec2:quota-volumetotalsize	Number of total volume size, in GB	account and user
iam:quota-groupnumber	Number of IAM groups	account
iam:quota-usernumber	Number of IAM users	account

Default Quota

Contrary to IAM policies, by default, there is no quota limits (except the hard system limit) on any resource allocations for a user or an account. Also, system administrators are not constrained by any quota. Account administrators are only be constrained by account quota.

Algorithms

Eucalyptus uses two types of algorithms to determine access.

Policy Evaluation Algorithm

You can associated multiple policies and permission statements with a user. The way these are combined together to control the access to resources in an account is defined by the policy evaluation algorithm. Eucalyptus implements the same policy evaluation algorithm as IAM:

1. If the request user is account admin, access is allowed.
2. Otherwise, collect all the policy statements associated with the request user (attached to the user and all the groups the user belongs to), which match the incoming request (i.e. based on the API being invoked and the resources it is going to access).
 - a. If there is no matched policy statement, access is denied (default implicit deny).
 - b. Otherwise, evaluate each policy statement that matches.
 - If there is a statement that explicitly denies the access, the request is denied.
 - If there is no explicit deny, which means there is at least one explicit allow, access is allowed.

Access Evaluation Algorithm

Now we give the overall access evaluation combining both account level permissions and IAM permissions, which decides whether a request is accepted by Eucalyptus:

1. If the request user is sys admin, access is allowed.
2. Otherwise, check account level permissions, e.g. image launch permission, to see if the request user's account has access to the specific resources.
 - a. If not, the access is denied.

- b. Otherwise, invoke the policy evaluation algorithm to check if the request user has access to the resources based on IAM policies.

Quota Evaluation Algorithm

Like the normal IAM policies, a user may be associated with multiple quota policies (and multiple quota statements). How all the quota policies are combined to take effect is defined by the quota evaluation algorithm:

1. If the request user is sys admin, there is no limit on resource usage.
2. Otherwise, collect all the quotas associated with the request user, including those attached to the request user's account and those attached to the request user himself/herself (for account admin, we only need collect account quotas).
3. Evaluate each quota one by one. Reject the request as long as there is one quota being exceeded by the request. Otherwise, accept the request.



Note: The hard limits on some resources override quota limits. For example, `walrus.storagemaxbucketsizeinmb` (system property) overrides the `s3:quota-bucketsize` (quota key).

Sample Policies

A few example use cases and associated policies.

Here are some example use cases and associated policies. You can edit these policies for your use, or use them as examples of JSON syntax and form.



Tip: For more information about JSON syntax used with AWS resources, go to [Using AWS Identity and Access Management](#).

Examples: Allowing Specific Actions

The following policy allows a user to only run instances and describe things.

```
{
  "Statement": [{
    "Effect": "Allow",
    "Action": [ "ec2:*Describe*", "ec2:*Run*" ],
    "Resource": "*"
  }]
}
```

The following policy allows a user to only list things:

```
{
  "Statement": [
    {
      "Sid": "Stmt1313686153864",
      "Action": [
        "iam:List*"
      ],
      "Effect": "Allow",
      "Resource": "*"
    }
  ]
}
```

The following policy grants a generic basic user permission for running instances and describing things.

```
{
```

```

"Statement": [
  {
    "Sid": "Stmt1313605116084",
    "Action": [
      "ec2:AllocateAddress",
      "ec2:AssociateAddress",
      "ec2:AttachVolume",
      "ec2:Authorize*",
      "ec2:CreateKeyPair",
      "ec2:CreateSecurityGroup",
      "ec2:CreateSnapshot",
      "ec2:CreateVolume",
      "ec2>DeleteKeyPair",
      "ec2>DeleteSecurityGroup",
      "ec2>DeleteSnapshot",
      "ec2>DeleteVolume",
      "ec2:Describe*",
      "ec2:DetachVolume",
      "ec2:DisassociateAddress",
      "ec2:GetConsoleOutput",
      "ec2:RunInstances",
      "ec2:TerminateInstances",
      "ec2:ReleaseAddress"
    ],
    "Effect": "Allow",
    "Resource": "*"
  }
]
}

```

Examples: Denying Specific Actions

The following policy allows a user to do anything but delete.

```

{
  "Statement": [
    {
      "Action": [
        "ec2>Delete*"
      ],
      "Effect": "Deny",
      "Resource": "*"
    }
  ]
}

```

The following policy denies a user from creating other users.

```

{
  "Statement": [
    {
      "Sid": "Stmt1313686153864",
      "Action": [
        "iam:CreateUser"
      ],
      "Effect": "Deny",
      "Resource": "*"
    }
  ]
}

```

Examples: Specifying Time Limits

The following policy allows a user to run instances within a specific time.

```
{
  "Statement": [
    {
      "Sid": "Stmt1313453084396",
      "Action": [
        "ec2:RunInstances"
      ],
      "Effect": "Allow",
      "Resource": "*",
      "Condition": {
        "DateLessThanEquals": {
          "aws:CurrentTime": "2011-08-16T00:00:00Z"
        }
      }
    }
  ]
}
```

The following policy blocks users from running instances at a specific time.

```
{
  "Statement": [
    {
      "Sid": "Stmt1313453084396",
      "Action": [
        "ec2:RunInstances"
      ],
      "Effect": "Allow",
      "Resource": "*",
      "Condition": {
        "DateLessThanEquals": {
          "aws:CurrentTime": "2011-08-16T00:00:00Z"
        }
      }
    }
  ]
}
```

The following policy keeps alive an instance for .

```
{
  "Statement": [
    {
      "Action": ["ec2:RunInstances" ],
      "Effect": "Allow",
      "Resource": "*",
      "Condition": { "NumericEquals":{"ec2:KeepAlive":"60000"}}
    }
  ]
}
```

The following policy sets an expiration date on running instances.

```
{
  "Statement": [
    {
      "Action": ["ec2:RunInstances" ],
      "Effect": "Allow",
```



```

    "Resource": "*",
    "Condition": { "DateEquals":
{"ec2:ExpirationTime": "2011-08-16T00:00:00Z"} }
  ]
}

```

Examples: Restricting Resources

The following policy allows users to only launch instances with a large image type.

```

{
  "Statement": [
    {
      "Action": [
        "ec2:RunInstances"
      ],
      "Effect": "Allow",
      "Resource": "arn:aws:ec2::vmtype/ml.xlarge"
    }
  ]
}

```

The following policy restricts users from launching instances with a specific image ID.

```

{
  "Statement": [
    {
      "Action": [
        "ec2:RunInstances"
      ],
      "Effect": "Deny",
      "Resource": "arn:aws:ec2::image/emi-0FFF1874"
    }
  ]
}

```

The following policy restricts users from allocating addresses to a specific elastic IP address.

```

{
  "Statement": [
    {
      "Sid": "Stmt1313626078249",
      "Action": "*",
      "Effect": "Deny",
      "Resource": "arn:aws:ec2::address/192.168.10.140"
    }
  ]
}

```

The following policy denies volume access.

```

{
  "Statement": [
    {
      "Action": [
        "ec2:*"
      ],
      "Effect": "Deny",
      "Resource": "arn:aws:ec2::volume/*"
    }
  ]
}

```

```
]
}
```

LDAP/AD Integration

You can use the Eucalyptus LDAP/Active Directory (AD) integration to synchronize existing LDAP/AD user and group information with Eucalyptus. When you enable LDAP/AD synchronization, Eucalyptus does the following:

- Imports specified user and group information from LDAP or AD and maps them into a predefined two-tier account/group/user structure
- Authenticates Dashboard users through the connected LDAP or AD service

Note that Eucalyptus only imports the identities and some related information. Any Eucalyptus-specific attributes are still managed from Eucalyptus. These include:

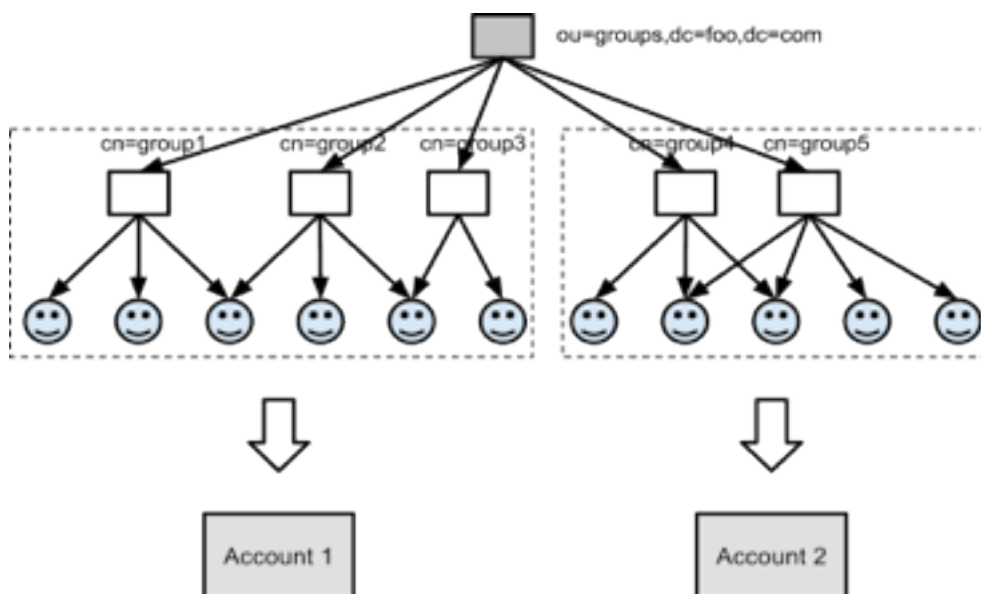
- User credentials: secret access keys and X.509 certificates. The Dashboard login password is an exception. Eucalyptus does not download passwords from LDAP/AD and does not save them either. Eucalyptus authenticates Dashboard logins directly through LDAP/AD, using LDAP/AD authentication (simple or SASL).
- Policies: IAM policies and quotas. Policies are associated with identities within Eucalyptus, and stored in internal database.

Also note that special identities, including system administrators and account administrators, are created in Eucalyptus and not imported from LDAP/AD. Only normal user identities are imported.

Identity Mapping

Identities in LDAP/AD are organized differently from the identity structure in Eucalyptus. So a transformation is required to map LDAP/AD identities into Eucalyptus.

The following image shows a simple scheme of how the mapping works. In this scheme, the user groups in LDAP tree are partitioned into two sets. Each set is mapped into one separate account. Group 1, 2 and 3 are mapped to Account 1 and Group 4 and 5 are mapped to Account 2. As the result, all users in Group 1, 2 and 3 will be in Account 1, and all users in Group 4 and 5 will be in Account 2.



To summarize the mapping method:

1. Pick user groups from LDAP/AD and combine them into different accounts. There are two ways of doing this:
 - Use something called accounting groups. Account groups are essentially groups of groups. Each accounting group contains multiple user groups in LDAP/AD. Then each accounting group maps to an account in Eucalyptus.

- Manually partition groups into accounts. Each group partition maps to an account.
2. Once the accounts are defined (by accounting groups or group partitions), all the LDAP/AD user groups will be mapped into Eucalyptus groups within specific accounts; and LDAP/AD users will be mapped into Eucalyptus users. You also have options to filter the groups and users to be imported into Eucalyptus.

Note that each group can be mapped into multiple accounts. But understand that Eucalyptus accounts are separate name spaces. So for groups and users that are mapped into different accounts, their information (name, attributes, etc) will be duplicated in different accounts. And duplicated users will have separate credentials in different accounts. For example, Group 1 may map to both Account 1 and Account 2. Say user A belongs to Group 1. Then Account 1 will have user A and Account 2 will also have user A. User A in Account 1 and user A in Account 2 will have different credentials, policies, etc., but the same user information.



Note: Currently, there is not a way to map individual users into an account. The mapping unit is LDAP user group.

LDAP/AD Integration Configuration

The LDAP/AD Integration Configuration (LIC) is a JSON format file. This file specifies everything Eucalyptus needs to know about how to synchronize with an LDAP or AD service.

You can find a LIC template at `/usr/share/eucalyptus/lic_template`. This template shows all the fields of the LIC, and provides detailed documentation and example values for each field.

To start a LIC file, use the LIC command line tool.

```
/usr/sbin/euca-lictool --password <password> --out example.lic
```

The above command invokes the LIC tool to create a template LIC and fill in the encrypted password for authenticating to LDAP/AD service (i.e. the password of the administrative user for accessing the LDAP/AD during synchronization). The LIC tool's primary functions are to encrypt the LDAP/AD password and to generate the starting LIC template. The usage of the LIC tool shows different ways to invoke the command.

Once you have the LIC template, you can fill in the details by editing the `*.lic` file using your favorite editor as it is a simple text file. As we said above, the LIC file is in JSON format. Each top level entity specifies one aspect of the LDAP/AD synchronization. The following shows one possible example of a LIC file.

```
{
  "ldap-service": {
    "server-url": "ldap://localhost:7733",
    "auth-method": "simple",
    "user-auth-method": "simple",
    "auth-principal": "cn=ldapadmin,dc=foo,dc=com",
    "auth-credentials": "{RSA/ECB/PKCS1Padding}EAXRnvwnKtCZOxSrD/F3ng/
yHH3J4jMxNUS
kJJf6oqNMsUihjUerZ20e5iyXImPgjK1ELAPnppEfJvhCs7woS7jtFsedunsp5DJCNhgmOb2CR/
MnH
11V3FNY7bBWoeW5A8Wwy6x7YrPMS0j7dJkwM7yfp1z6AbKOo2688I9uIvJUQwEKS4dOp7RVdA0iz1J
BDPAXiFZ2qa40VjFI/1mggbiWDNlgxiVtZXAEK7x9SRHJytLS8nrNPPiVpuTg3djKiWPVOLZ6vpSgP
cVeliP261qdUfnf3GDKi3jqbPpRRQ6n8yI6aHw0gAtq8/
qPyqjkkDP8JsGBgmXMxiCNPogbWg==",
    "use-ssl": "false",
    "ignore-ssl-cert-validation": "false",
    "krb5-conf": "/path/to/krb5.conf",
  },
  "sync": {
    "enable": "true",
    "auto": "true",
  }
}
```

```

    "interval": "900000",
    "clean-deletion": "false",
  },
  "accounting-groups": {
    "base-dn": "ou=groups,dc=foo,dc=com",
    "id-attribute": "cn",
    "member-attribute": "member",
    "selection": {
      "filter": "objectClass=accountingGroup",
      "select": [ "cn=accountingToSelect,ou=Groups,dc=foo,dc=com" ],
      "not-select": [ "cn=accountingToIgnore,ou=Groups,dc=foo,dc=com" ],
    }
  },
  "groups": {
    "base-dn": "ou=groups,dc=foo,dc=com",
    "id-attribute": "cn",
    "member-attribute": "member",
    "selection": {
      "filter": "objectClass=groupOfNames",
      "select": [ "cn=groupToSelect,ou=Groups,dc=foo,dc=com" ],
      "not-select": [ "cn=groupToIgnore,ou=Groups,dc=foo,dc=com" ],
    }
  },
  "users": {
    "base-dn": "ou=people,dc=foo,dc=com",
    "id-attribute": "uid",
    "user-info-attributes": {
      "fullName": "Full Name",
      "email": "Email"
    },
    "selection": {
      "filter": "objectClass=inetOrgPerson",
      "select": [ "uid=john,ou=People,dc=foo,dc=com",
        "uid=jack,ou=People,dc=foo,dc=com" ],
      "not-select": [ "uid=tom,ou=People,dc=foo,dc=com" ],
    }
  },
},

```

In the following sections explain each field of LIC in detail.

ldap-service

The `ldap-service` element contains everything related to the LDAP/AD service.

Element	Description
server-url	The LDAP/AD server URL, starting with <code>ldap://</code> or <code>ldaps://</code> .
auth-method	The LDAP/AD authentication method to perform synchronization.
auth-principal	The ID of the administrative user for synchronization.
auth-credentials	The credentials for LDAP/AD authentication, like a password. We recommend that you encrypt this using <code>usr/sbin/euca-lictool</code> .
user-auth-method	The LDAP/AD authentication method for normal users to perform Dashboard login. <ul style="list-style-type: none"> <i>simple</i>: for clear text user/password authentication.

Element	Description
use-ssl	<ul style="list-style-type: none"> <i>DIGEST-MD5</i>: for SASL authentication using MD5 <i>GSSAPI</i>: SASL authentication using Kerberos V5.
ignore-ssl-cert-validation	Specifies whether to use SSL for connecting to LDAP/AD service.
krb5-conf	Specifies whether to ignore self-signed SSL certs. This is useful when you only have self-signed SSL certs for your LDAP/AD services.
	The file path for krb5.conf, if you use GSSAPI authentication method.

sync

The `sync` element contains elements for controlling synchronization.

Element	Description
enable	States whether the Dashboard uses LDAP rather than the Eucalyptus database for log-ins. Set to true to turn on LDAP web logging. Set to false to use the Eucalyptus database for web logging. If set to false, you can ignore all other fields in this section. Default value: false
auto	Set to true to turn on automatic synchronization. Set to false to turn off synchronization.
interval	The length in milliseconds of the automatic synchronization interval.
clean-deletion	Parameter denoting whether to remove identity entities from Eucalyptus when they are deleted from LDAP. Set to true if you want Eucalyptus to remove any identities once their counterparts in LDAP are deleted. Set to false if you want these identities kept without being purged.

accounting-groups

This section uses a special group in LDAP/AD to designate accounts in the Eucalyptus “accounting group.” The accounting group takes normal LDAP/AD groups as members, i.e., they are groups of groups. The accounting group’s name becomes the account name in Eucalyptus. The member groups become Eucalyptus groups in that account. And the users of all those groups become Eucalyptus users within that account and corresponding Eucalyptus groups.



Important: If you use `accounting-groups`, remove the `groups-partition` section. These two sections are mutually exclusive.

Element	Description
base-dn	The base DN of accounting groups in the LDAP/AD tree.
id-attribute	The ID attribute name of the accounting group entry in LDAP/AD tree.
member-attribute	The LDAP/AD attribute name for members of the accounting group.
selection	The accounting groups you want to map to. This contains the following elements:

Element	Description
	<ul style="list-style-type: none"> <i>filter</i>: The LDAP/AD searching filter used for the LDAP/AD search to get the relevant LDAP/AD entities, e.g. the users to be synchronized. (Example: <code>objectClass=groupOfNames</code>) <i>select</i>: Explicitly gives the full DN of entities to be synchronized, in case they can not be specified by the search filter. (Example: <code>cn=groupToSelect,ou=Groups,dc=foo,dc=com</code>) <i>not-select</i>: Explicitly gives the full DN of entities NOT to be synchronized, in case this can not be specified by the search filter. (Example: <code>cn=groupToIgnore,ou=Groups,dc=foo,dc=com</code>)

groups-partition

Like accounting-groups, groups-partition specifies how to map LDAP/AD groups to Eucalyptus accounts. However, in this section you to manually specify which LDAP/AD groups you want to map to Eucalyptus accounts.



Important: If you use groups-partition, remove the accounting-groups section. These two sections are mutually exclusive.

The Eucalyptus accounts are created by partitioning LDAP/AD groups. Each partition composes an Eucalyptus account. So all the groups within the partition become Eucalyptus groups within that account. All the users of those groups will become Eucalyptus users within that account and the corresponding Eucalyptus groups.

This section requires that you specify one partition at a time, using a list of JSON key-value pairs. For each entry, the key is the account name to be mapped and the value is a list of names of LDAP/AD groups to be mapped into the account. For example:

```
"groups-partition": {
  "salesmarketing": ["sales", "marketing"],
  "devsupport": ["engineering", "support"],
}
```

Here salesmarketing and devsupport are names for the groups partition and are used as the corresponding Eucalyptus account names.



Tip: If you use groups-partition, remove the accounting-groups section. These two sections are mutually exclusive.

groups

Short reference description.

This groups element specifies how to map LDAP/AD groups to Eucalyptus groups. It contains the elements listed in the following table. The meanings are similar to those in accounting-groups element.

Element	Description
base-dn	The base DN for searching groups.
id-attribute	The ID attribute name of the LDAP group.
member-attribute	The name of the attribute for group members. Usually, it is member in modern LDAP implementation, which lists full user DN.
selection	The specific LDAP/AD groups you want to map to. This contains the following elements:

Element	Description
	<ul style="list-style-type: none"> <i>filter</i>: The LDAP/AD searching filter used for the LDAP/AD search to get the relevant LDAP/AD entities, e.g. the users to be synchronized. (Example: objectClass=groupOfNames) <i>select</i>: The LDAP/AD searching filter used for the LDAP/AD search to get the relevant LDAP/AD entities, e.g. the users to be synchronized. (Example: objectClass=groupOfNames) <i>not-select</i>: Explicitly gives the full DN of entities NOT to be synchronized, in case this can not be specified by the search filter. (Example: cn=groupToIgnore,ou=Groups,dc=foo,dc=com)

users

Explicitly gives the full DN of entities NOT to be synchronized, in case this can not be specified by the search filter. (Example: cn=groupToIgnore,ou=Groups,dc=foo,dc=com)

Element	Description
base-dn	The base DN for searching users.
id-attribute	The attribute ID of the LDAP user.
selection	<p>The specific LDAP/AD users you want to map to. This contains the following elements:</p> <ul style="list-style-type: none"> <i>filter</i>: The LDAP/AD searching filter used for the LDAP/AD search to get the relevant LDAP/AD entities, e.g. the users to be synchronized. (Example: objectClass=groupOfNames) <i>select</i>: Explicitly gives the full DN of entities to be synchronized, in case they can not be specified by the search filter. (Example: cn=groupToSelect,ou=Groups,dc=foo,dc=com) <i>not-select</i>: Explicitly gives the full DN of entities NOT to be synchronized, in case this can not be specified by the search filter. (Example: cn=groupToIgnore,ou=Groups,dc=foo,dc=com)

Synchronization Process

The synchronization always starts when the following happens:

- You manually upload a LDAP/AD Integration Configuration (LIC) file. Every new or updated LIC upload triggers a new synchronization.
- If the automatic synchronization is enabled, a synchronization is started when the timer goes off.



Note: Eucalyptus does not allow concurrent synchronization. If you trigger synchronization more than once within a short time period, Eucalyptus only allows the first one.

During a synchronization, everything specified by an LIC in the LDAP/AD tree will be downloaded into Eucalyptus' internal database. Each synchronization is a merging process of the information already in the database and the information from LDAP/AD. There are three cases for each entity: user, group or account:

- If an entity from LDAP/AD is not in Eucalyptus, a new one is created in the database.

- If an entity from LDAP/AD is already in Eucalyptus, the Eucalyptus version is updated. For example, if a user's info attributes are changed, those changes are downloaded and updated.
- If an entity in Eucalyptus is missing from LDAP/AD, it will be removed from the database if the clean-deletion option in LIC is set to true. Otherwise, it will be left in the database.



Important: If clean-deletion is set to true, the removed entities in Eucalyptus will be lost forever, along with all its permissions and credentials. The resources associated with the entity will be left untouched. It is system administrator's job to recycle these resources.

Access Tasks

This section provides details about the tasks you perform using policies and identities. The tasks you can perform are divided up into tasks for users, tasks for groups, and tasks for policies.

The following use cases detail work flows for common processes:

- [*Use Case: Creating an Administrator*](#)
- [*Use Case: Creating a User*](#)

You can perform the following access-related tasks listed in the following sections:

- Accounts:
 - [*Add an Account*](#)
 - [*Approve an Account*](#)
 - [*Reject an Account*](#)
 - [*Rename an Account*](#)
 - [*List Accounts*](#)
 - [*Delete an Account*](#)
- Groups:
 - [*Create a Group*](#)
 - [*Add a Group Policy*](#)
 - [*Modify a Group*](#)
 - [*Add a User to a Group*](#)
 - [*Remove a User from a Group*](#)
 - [*List Groups*](#)
 - [*List Policies for a Group*](#)
 - [*Delete a Group*](#)
- Users:
 - [*Add a User*](#)
 - [*Create a Login Profile*](#)
 - [*Modify a User*](#)
 - [*List Users*](#)
 - [*Delete a User*](#)
- Credentials:
 - [*Create Credentials*](#)
 - [*Get Credentials*](#)
 - [*Generate Administrator Credentials*](#)
 - [*Upload a Certificate*](#)

Use Case: Creating an Administrator

This use case details tasks for creating an administrator. These tasks require that you have your account credentials for sending requests to Eucalyptus using the command line interface (CLI) tools.

Create an Admin Group

Eucalyptus recommends using account credentials as little as possible. You can avoid using account credentials by creating a group of users with administrative privileges.

1. Create a group called administrators.

```
euare-groupcreate -g administrators
```

2. Verify that the group was created.

```
euare-grouplistbypath
```

Eucalyptus returns a listing of the groups that have been created, as in the following example.

```
arn:aws:iam::123456789012:group/administrators
```

Add a Policy to the Group

Add a policy to the administrators group that allows its members to perform all actions in Eucalyptus.

Enter the following command to create a policy called admin-root that grants all actions on all resources to all users in the administrators group:

```
euare-groupaddpolicy -p admin-root -g administrators -e Allow -a "*" -r
"*" -o
```

Create an Administrative User

Create a user for day-to-day administrative work and add that user to the administrators group.

1. Enter the following command to create an administrative user named alice:

```
euare-usercreate -u alice
```

2. Add the new administrative user to the administrators group.

```
euare-groupadduser -g administrators -u alice
```

Generate Administrative Credentials

To start running commands as the new administrative user, you must create an access key for that user.

1. Enter the following command to generate an access key for the administrative user:

```
euare-useraddkey -u alice
```

Eucalyptus returns the access key ID and the user's secret key.

2. Open the ~/.eucarc file and replace your account credentials you just created, as in this example:

```
export EC2_ACCESS_KEY='WOKSEQRNM1LVIR702XVX1 '
export EC2_SECRET_KEY='0SmLCQ8DAZPKoaC7oJYcRMfeDUgGbiSVv1ip5WaH '
```

3. Save and close the file.

4. Open the `~/.iamrc` file and replace your account credentials, as in this example:

```
AWSAccessKeyId=WOKSEQRNM1LVIR702XVX1
AWSSecretKey=0SmLCQ8DAZPKoaC7oJYcRMfeDUgGbiSVv1ip5WaH
```

5. Save and close the file.
6. Switch `euca2ools` over to using the new credentials.

```
source ~/.eucarc
```

Use Case: Creating a User

This use case details tasks needed to create a user with limited access.

Create a Group

We recommend that you apply permissions to groups, not users. In this example, we will create a group for users with limited access.

1. Enter the following command to create a group for users who will be allowed create snapshots of volumes in Eucalyptus.

```
euare-groupcreate -g ebs-backup
```

2. Open an editor and enter the following JSON policy:

```
{
  "Statement": [
    {
      "Action": [
        "ec2:CreateSnapshot"
      ],
      "Effect": "Allow",
      "Resource": "*"
    }
  ]
}
```

3. Save and close the file.
4. Enter the following to add the new policy name `allow-snapshot` and the JSON policy file to the `ebs-backup` group:

```
euare-groupuploadpolicy -g ebs-backup -p allow-snapshot -f allow-snapshot.json
```

Create the User

Create the user for the group with limited access.

Enter the following command to create the user `sam` in the group `ebs-backup` and generate a new key pair for the user:

```
euare-usercreate -u sam -g ebs-backup -k
```

Eucalyptus responds with the access key ID and the secret key, as in the following example:

```
AKIAJ25S6IJ5K53Y5GCA
QLKyiCpfjWAVlo9pWqWCbuGB9L3T61w7nYYF0571
```

Accounts

Accounts are the primary unit for resource usage accounting. Each account is a separate name space and is identified by its UUID (Universal Unique Identifier).

Tasks performed at the account level can only be done by the users in the **eucalyptus** account.

Add an Account

To add a new account:

Add an Account (CLI)

To add a new account using the CLI:

Enter the following command:

```
euare-accountcreate -a <account_name>
```

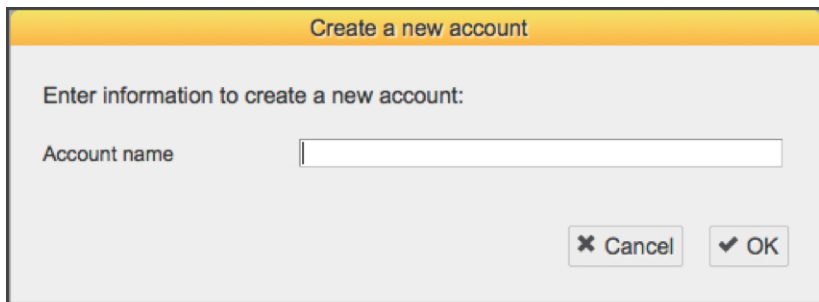
Eucalyptus returns the account name and its ID, as in this example:

```
account01 592459037010
```

Add an Account (Dashboard)

To add a new account using the Dashboard:

1. Click **Accounts** in the Quick Links section.
The **Accounts** page displays.
2. Click **New account** in the **Accounts** page.
The **Create a new account** popup displays.
3. Enter an account name in the **Account name** field and click **OK**.



The new account displays in the list on the **Accounts** page.

Approve an Account

To approve an account:

1. Click **Accounts** in the Quick Links section.
The **Accounts** page displays.
2. Click the **ID** of the account you want to approve.
The account, name, and Registration status are highlighted.
3. Click **Approve**.
The **Approve selected accounts** popup displays.
4. Verify that the displayed account is the one you want, and click **OK**.

The account's registration status displays as **CONFIRMED** on the **Accounts** page.

Reject an Account

To reject an account:

1. Click **Accounts** in the Quick Links section.
The **Accounts** page displays.
2. Click the **ID** of the account you want to delete.
The account, name, and Registration status are highlighted.
3. Click **Reject**.
The **Reject selected accounts** popup displays.
4. Verify that the displayed account is the one you want, and click **OK**.

The account no longer displays on the **Accounts** page.

Rename an Account

This section explains steps to perform so that you can rename an account.

Using the CLI

To change an account's name using the CLI:

Enter the following command:

```
euare-accountaliascreate -a <new_name>
```

Using the Dashboard

To change an account's name using the Dashboard:

1. Click **Accounts** in the Quick Links section.
The **Accounts** page displays.
2. Click the **ID** of the account you want to rename.
The account's **Properties** area displays.
3. In the **Name** field, enter the new name of the account.
4. Click **Save**.

The new account name displays in the **Accounts** page.

List Accounts

This section explains steps to perform so that you can list accounts.

Using the CLI

Use the `euare-accountlist` command to list all the accounts in an account or to list all the users with a particular path prefix. The output lists the ARN for each resulting user.

```
euare-userlistbypath -p <path>
```

Using the Dashboard

To list accounts using the Dashboard:

Click **Accounts** in the Quick Links section.
The **Accounts** page displays.

The **Accounts** page displays all accounts in your cloud.

Delete an Account



Tip: If there are resources tied to the account that you delete, the resources remain. We recommend that you delete these resources first.

To delete an account:

Delete an Account (CLI)

To delete an account using the CLI:

Enter the following command:

```
euare-accountdel -a <account_name> -r true
```

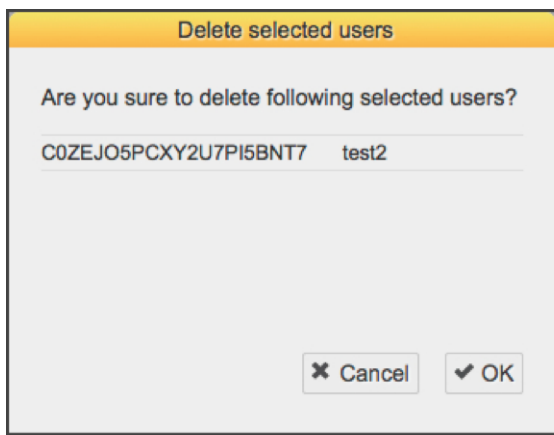
Use the `-r` option set to `true` to delete the account recursively. You don't have to use this option if have already deleted users, keys, and passwords in this account.

Eucalyptus does not return any message.

Delete an Account (Dashboard)

To delete an account:

1. Click **Accounts** in the Quick Links section.
The **Accounts** page displays.
2. Click the **ID** of the account you want to delete.
The account, name, and Registration status are highlighted.
3. Click **Delete accounts**.
The **Delete selected accounts** popup displays.
4. Verify that the displayed account is the one you want, and click **OK**.



The account no longer displays in the list on the **Accounts** page.

Groups

Groups are used to share resource access authorizations among a set of users within an account. Users can belong to multiple groups.



Important: A group in the context of access is not the same as a security group.

This section details tasks that can be performed on groups.

Create a Group

Using the CLI

To create a group using the CLI:

Enter the following command:

```
euare-groupcreate -g <group_name>
```

Eucalyptus does not return anything.

Using the Dashboard

To create a group using the Dashboard:

1. Click **Accounts** in the Quick Links section.
The **Accounts** page displays.
2. Click the **ID** of the account you want to add a group to.
The account, name, and Registration status are highlighted.
3. Click **New groups** in the **Accounts** page.
The **Create new groups** popup displays.
4. Enter the group name in the **Group name** field.



Tip: You can add more than one group at a time. Every group you add, however, will be in the same path.

5. Enter the group path in the **Group path** field.
6. Click **OK**.

The group is associated with the account you chose. You can see the information if you select the account in the **Accounts** page and click the **Member groups** link, located in the **Properties** section of the screen.

Add a Group Policy

Using the CLI

To add a policy to a group using the CLI:

Enter the following command:

```
euare-groupaddpolicy -g <group_name> -p <policy_name> -e <effect> -a  
  <actions> -o
```

The optional `-o` parameter tells Eucalyptus to return the JSON policy, as in this example:

```
{ "Version": "2008-10-17", "Statement": [ { "Effect": "Allow", "Action":  
  [ "ec2:RunInstances" ], "Resource": [ "*" ] } ] }
```

Using the Dashboard

To add a policy to a group using the Dashboard:

1. Click **Groups** in the Quick Links section.
The **Groups** page displays.
2. Click the **ID** of the group you want to add a policy to.
The ID, Name, Path, and Owner account line is highlighted.
3. Click **Add policy**.
The **Add new policy** popup displays.

4. Enter the policy name in the **Policy name** field.
5. Enter the policy content in the **Policy content** field.
6. Click **OK**.

The policy is now added to the group.

Modify a Group

Modifying a group is similar to a "move" operation. Whoever wants to modify the group must have permission to do it on both sides of the move. That is, you need permission to remove the group from its current path or name, and put that group in the new path or name.

For example, if a group changes from one area in a company to another, you can change the group's path from /area_abc/ to /area_efg/. You need permission to remove the group from /area_abc/. You also need permission to put the group into /area_efg/. This means you need permission to call `UpdateGroup` on both `arn:aws:iam::123456789012:group/area_abc/*` and `arn:aws:iam::123456789012:group/area_efg/*`.

Using the CLI

To modify a group using the CLI:

1. Enter the following command to modify the group's name:

```
euare-groupmod -g <group_name> --new-group-name <new_name>
```

Eucalyptus does not return a message.

2. Enter the following command to modify a group's path:

```
euare-groupmod -g <group_name> -p <new_path>
```

Eucalyptus does not return a message.

Using the Dashboard

To modify a group using the Dashboard:

1. Click **Groups** in the Quick Links section.
The **Groups** page displays.
2. Click the **ID** of the group you want to rename.
The group's **Properties** area displays.
3. In the **Name** field, enter the new name of the group.
4. In the **Path** field, enter the new path for the group.
5. Click **Save**.

The new group name displays in the **Groups** page.

Remove a User from a Group

Using the CLI

To remove a user from a group using the CLI:

Enter the following command:

```
euare-groupremoveuser -g <group_name> -u <user-name>
```

Using the Dashboard

To remove a user from a group using the Dashboard:

1. Click **Groups** in the Quick Links section.

The **Groups** page displays.

2. Click the **ID** of the group you want to remove the user from.
The ID, Name, Path, and Owner account line is highlighted.
3. Click **Remove users**.
The **Remove users to selected groups** popup displays.

4. Enter the name of the user you want to remove and click **OK**.

The user is now removed from the group.

List Groups Using the CLI

To list all the groups a specific user is in:

Enter the following command:

```
euare-grouplistbypath
```

Eucalyptus returns a list of paths followed by the ARNs for the groups in each path. For example:

```
arn:aws:iam::eucalyptus:group/groupa
```

Using the Dashboard

To list groups using the Dashboard:

Click **Groups** in the Quick Links section.

The **Groups** page displays.

The **Groups** page displays all groups in your cloud.

List Policies for a Group

Using the CLI

To list policies associated with a group using the CLI:

Enter the following command:

```
euare-grouplistpolicies -g <group_name>
```

Eucalyptus returns a listing of all policies associated with the group.

Using the Dashboard

To list the policies associated with a group using the Dashboard:

1. Click **Groups** in the Quick Links section.
The **Groups** page displays.
2. Click the **ID** of the group you want to list policies for.
The **Properties** section displays.
3. In the **Properties** section, click **Policies**.
The **Access Policies** page displays.

Delete a Group

Using the CLI

When you delete a group, you have to remove users from the group and delete any policies from the group. You can do this with one command, using the `euare-groupdel` command with the `-r` option. Or you can follow the following steps to specify who and what you want to delete.

1. Individually remove all users from the group.

```
euare-groupremoveuser -g <group_name> -u <user_name>
```

2. Delete the policies attached to the group.

```
euare-groupdelpolicy -g <group_name> -p <policy_name>
```

3. Delete the group.

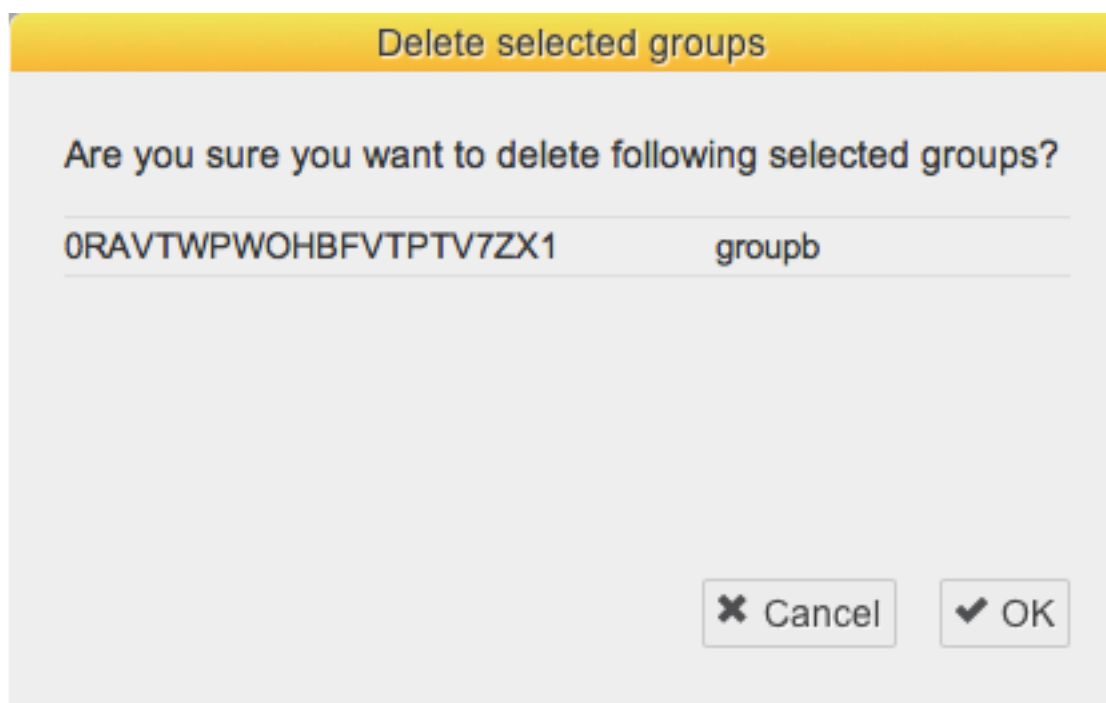
```
euare-groupdel -g <group_name>
```

The group is now deleted.

Using the Dashboard

To delete a group using the Dashboard:

1. Click **Groups** in the Quick Links section.
The **Groups** page displays.
2. Click the **ID** of the group you want to delete.
The ID, Name, Path, and Owner account line is highlighted.
3. Click **Delete groups**.
The **Delete selected groups** popup displays.



4. Click **OK**.

The group is now deleted.

Users

Users are subsets of accounts and are added to accounts by an appropriately credentialed administrator. While the term **user** typically refers to a specific person, in Eucalyptus, a **user** is defined by a specific set of credentials generated to enable access to a given account. Each set of user credentials is valid for accessing only the account for which they were created. Thus a user only has access to one account within a Eucalyptus system. If an individual person wishes to have access to more than one account within a Eucalyptus system, a separate set of credentials must be generated (in effect a new 'user') for each account (though the same username and password can be used for different accounts).

When you need to add a new user to your Eucalyptus cloud, you'll go through the following process:

1	<i>Create a user</i>
2	<i>Add user to a group</i>
3	<i>Give user a login profile</i>

Add a User

Using the CLI

To add a user using the CLI:

Enter the following command

```
euare-usercreate -u <user_name> -g <group_name> -k
```

Eucalyptus does not return a response.



Tip: If you include the `-v` parameter, Eucalyptus returns a response that includes the user's ARN and GUID.

Using the Dashboard

To add a user using the Dashboard:

1. Click **Accounts** in the Quick Links section.
The **Accounts** page displays.
2. Click the **ID** of the account you want to rename.
The account's **Properties** area displays.
3. Click **New Users**.
The **Create new users** popup window displays.
4. Enter a name in the **User names** field.



Tip: You can add more than one user at a time. Every user you add, however, will be in the same path.

5. Enter a path in the **User path** field.
6. Click **OK**.

The user is added to the account.

Add a User to a Group

Using the CLI

To add a user to a group using the CLI:

Enter the following command:

```
euare-groupadduser -g <group_name> -u <user-name>
```

Using the Dashboard

1. Click **Groups** in the Quick Links section.
The **Groups** page displays.
2. Click the **ID** of the group you want to add the user to.
The ID, Name, Path, and Owner account line is highlighted.
3. Click **Add users**.
The **Add users to selected groups** popup displays.

4. Enter the name of the user you want to add and click **OK**.

The user is now added to the group.

Create a Login Profile

Once you create a user, you must generate a password for the user to use the Dashboard.

Using the CLI

To create a login profile using the CLI:

Enter the following command:

```
euare-useraddloginprofile -u <user_name> -p <password>
```

Eucalyptus does not return a response.

Using the Dashboard

To create a login profile using the Dashboard:

1. Click **Users** in the Quick Links section.
The **Users** page displays.
2. Click the **ID** of the user whose path you want to change.
The user's **Properties** area displays.
3. Click **Password**.
The **Change password** popup window displays.
4. Enter the new password in the **New user password** field, and repeat in the **New password again** field.
5. Click **OK**.

The login profile is now complete. If you are generating the password for a different user, let the user know the password and the URL to the Dashboard.

Modify a User

Modifying a user is similar to a "move" operation. Whoever wants to modify a user must have permission to do it on both sides of the move. That is, you need permission to remove the user from the current path or name, and put that user in the new path or name.

For example, if a user changes from one team in a company to another, you can change the user's path from `/team_abc/` to `/team_efg/`. You need permission to remove the user from `/team_abc/`. You also need permission to put the user into `/team_efg/`. This means you need permission to call `UpdateUser` on both `arn:aws:iam::123456789012:user/team_abc/*` and `arn:aws:iam::123456789012:user/team_efg/*`.

Using the CLI

To rename a user using the CLI:

1. Enter the following command to rename a user:

```
euare-usermod -u <user_name> --new-user-name <new_name>
```

Eucalyptus does not return a message.

2. Enter the following command:

```
euare-groupmod -u <user_name> -p <new_path>
```

Eucalyptus does not return a message.

Using the Dashboard

To rename a user using the Dashboard:

1. Click **Users** in the Quick Links section.
The **Users** page displays.

2. Click the **ID** of the user you want to rename.
The user's **Properties** area displays.
3. In the **Name** field, enter the new name of the user.
4. Click **Save**.

The new user name displays in the **Users** page.

List Users

You can list users within a path.

Using the CLI

Use the `euare-userlistbypath` command to list all the users in an account or to list all the users with a particular path prefix. The output lists the ARN for each resulting user.

```
euare-userlistbypath -p <path>
```

Using the Dashboard

To list users in the same path using the Dashboard:

1. Click **Users** in the Quick Links section.
The **Users** page displays.
2. Click the **Path** column to sort all users by path.

Delete a User

Using the CLI

To delete a user using the CLI:

Enter the following command

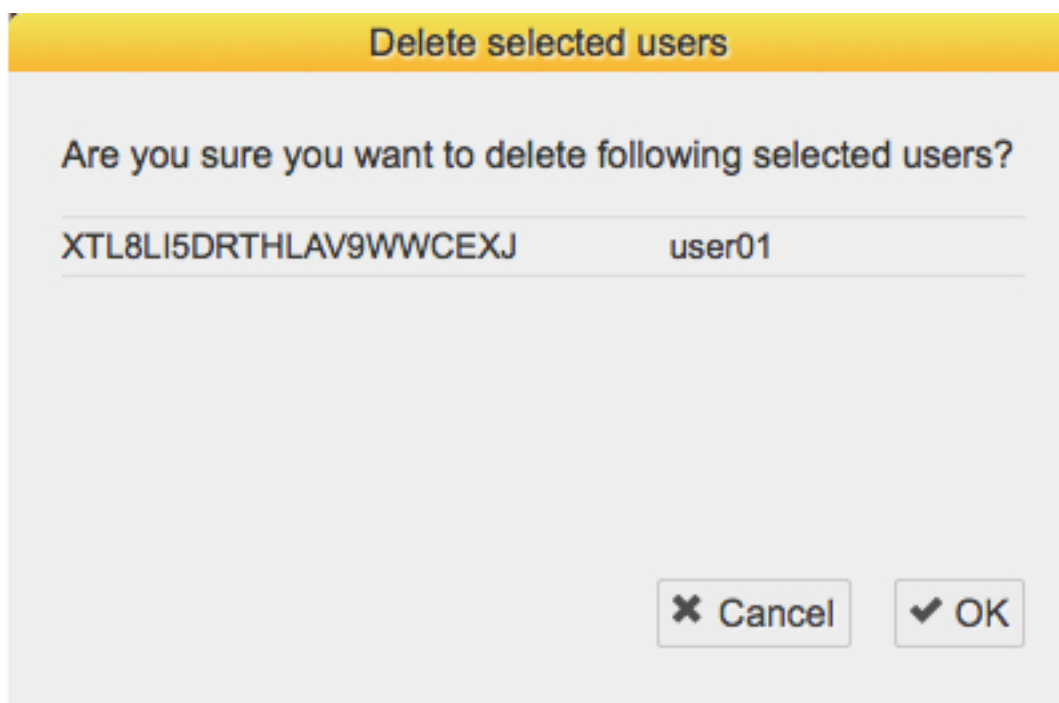
```
euare-userdel -u <user_name>
```

Eucalyptus does not return a response.

Using the Dashboard

To delete a user using the Dashboard:

1. Click **Users** in the Quick Links section.
The **Users** page displays.
2. Click the **ID** of the user you want to delete.
The user's information is highlighted.
3. Click **Delete Users**.
The **Delete selected users** popup window displays.



4. Click **OK**.

The user is deleted.

Credentials

Eucalyptus uses different types of credentials for different purposes. This section details tasks needed to allow access to Eucalyptus services.

Create Credentials

You can generate new credentials a number of ways. The first time you get credentials using either the Dashboard or the `euca_conf` command, a new secret access key is generated. On each subsequent request to get credentials, an existing active secret Key is returned. You can also generate new keys using the `euca-useraddkey` command.



Tip: Each request to get a user's credentials either via the download link in the Dashboard or using `euca_conf`, a new pair of a private key and X.509 certificate

- To generate a new key for a user by an account administrator, enter the following

```
euare-useraddkey -u <user_name>
```

- To generate a private key and an X.509 certificate pair, enter the following:

```
euare-usercreatecert -u <user_name>
```

Get Credentials

Eucalyptus provides two main ways of getting user credentials. In both cases, Eucalyptus returns a zip file that contains keys, certificates, a bash script, and several other required files. To use these credentials with such CLI tools as `euca2ools` or `ec2-tools`, unzip your credentials zip file to a directory of your choice.

- An administrator with a root access to the machine on which CLC is installed can get credentials using euca_conf CLI tool on that machine.

```
/usr/sbin/euca_conf --cred-account <account> --cred-user <user_name>
--get-credentials <filename>.zip
```

Where <account> and <user_name> are the names of the account and the user whose credentials are retrieved.



Tip: You can omit the --cred-account and --cred-user options when you get credentials for the **admin** user of the **eucalyptus** account.

- A user can get his or her credentials by logging in into the Dashboard and clicking **Download new credentials** in the drop-down menu at the top of the screen. This will result in a download of a zip file.

In the following example we download the credentials zip file to ~/ .euca, then change access permissions, as shown:

```
mkdir ~/.euca
cd ~/.euca
unzip <filepath>/<creds_zipfile>.zip
chmod 0700 ~/.euca
chmod 0600 *
```



Important: The zip file with credentials contains security-sensitive information. We recommend that you remove or read- and write-protect the file from other users after unzipping.

Alternatively, you can view and copy your access keys and X.509 certificates from the Dashboard after logging in, using the Navigation menu.

Generate Administrator Credentials



Important: When you run the following command, you are requesting a new X.509 and a corresponding private key. You cannot retrieve an existing private key.

To generate a set of credentials:

1. Log in to the CLC.
2. Get administrator credentials and source eucarc:

```
/usr/sbin/euca_conf --get-credentials admin.zip
unzip admin.zip
chmod 0600 *
source eucarc
```

Upload a Certificate

To upload a certificate provided by a user:

Enter the following command:

```
euare-useraddcert -u <user_name> -f <cert_file>
```

Synchronize LDAP/AD

To start an LDAP/AD synchronization:

1. Create an LDAP/AD Integration Configuration (LIC) file to specify all the details about the LDAP/AD synchronization.
2. Upload the LIC file to Eucalyptus using `euca-modify-property`.

Start a LIC File

The LIC is a file in JSON format, specifying everything Eucalyptus needs to know about how to synchronize with an LDAP or AD service. Eucalyptus provides a LIC template at `${EUCALYPTUS}/usr/share/eucalyptus/lic_template`. This template shows all the fields of the LIC, and provides detailed documentation and example values for each field.

To start a LIC file:

Enter the following command:

```
/usr/sbin/euca-lictool --password secret --out example.lic
```

The above command invokes the LIC tool to create a template LIC and fill in the encrypted password for authenticating to LDAP/AD service (i.e. the password of the administrative user for accessing the LDAP/AD during synchronization). The LIC tool's primary functions are to encrypt the LDAP/AD password and to generate the starting LIC template. The usage of the LIC tool shows different ways to invoke the command.

Once you have the LIC template, you can fill in the details by editing the `*.lic` file using a text editor. Each top level entity specifies one aspect of the LDAP/AD synchronization.

Upload a New LIC File

To upload a new LIC file:

Enter the following:

```
/usr/sbin/euca-modify-property -f  
authentication.ldap_integration_configuration=<lic_filename.lic>
```

This triggers a new synchronization using the uploaded LIC file.

Managing Images

A Eucalyptus Machine Image (EMI) is a special type of pre-packaged operating system and application software that Eucalyptus uses to create a virtual machine instance. An EMI contains all necessary information to boot instances of your software. An EMI can contain whatever software you want to meet your needs. For example, you can create an EMI to act as an application server, or a web server, or a custom web service. The following sections contain information to help you create, configure, and modify EMIs.



Important: All users can upload and register images, depending on their access. However, only the admin user can upload and register kernels or ramdisks.

Distro	Version	Arch	Hypervisor
CentOS	5	32	Xen
CentOS	5	32	esx-4.0
CentOS	5	32	esx-4.1
CentOS	5	32	esxi-4.0
CentOS	5	32	esxi-4.1
CentOS	5	64	Xen
CentOS	5	64	esx-4.0
CentOS	5	64	esx-4.1
CentOS	5	64	esxi-4.0
CentOS	5	64	esxi-4.1
RHEL	5	32	Xen
RHEL	5	64	Xen
RHEL	5	64	esx-4.0
RHEL	5	64	esx-4.1
RHEL	5	64	esxi-4.0
RHEL	5	64	esxi-4.1
RHEL	6	64	KVM
RHEL	6	64	esx-4.0
RHEL	6	64	esx-4.1
RHEL	6	64	esxi-4.0
RHEL	6	64	esxi-4.1
Ubuntu	Lucid	32	KVM
Ubuntu	Lucid	64	KVM
Ubuntu	Lucid	64	esx-4.0
Ubuntu	Lucid	64	esx-4.1
Ubuntu	Lucid	64	esxi-4.0
Ubuntu	Lucid	64	esxi-4.1

Image Overview

An image defines what will run on a guest instance with your Eucalyptus cloud. Typically, an image contains one of the Linux distributions like CentOS, Fedora, Ubuntu, Debian or others. It could also contain one of the supported Windows server versions. The format for these is identical.

Normally when we use the term "image" we mean the root file system. Once bundled, uploaded, and registered with Eucalyptus, such an image is known as a Eucalyptus machine image (EMI).

There are, however, other types of images that support the EMI. They are the kernel (EKI) and ramdisk (ERI). They contain kernel modules necessary for proper functioning of the image. Often, one set of these ERIs and EKIs are used by multiple EMIs. Once loaded into the Eucalyptus cloud, the EKI and ERI are referred to by the image and you don't have much interaction with them directly.



Tip: When you run an image, you can override the image's associated kernel and ramdisk if necessary (for example, if you want to try out another kernel. For more information, see [Associate a Kernel and Ramdisk](#).

To help get you started, Eucalyptus provides pre-packaged virtual machines that are ready to run in your cloud. You can get them at the [Eucalyptus Machine Images](#) page. Each Eucalyptus image from this site comes bundled with a corresponding EKI and ERI.

If you find that the pre-packaged images don't meet your needs, you can add an image from an existing, non-registered image, or create your own image.

Once you've selected and downloaded the image(s) you plan to use, read the details in the following section for directions about how to bundle, upload and register the images with your Eucalyptus cloud.

To view the status of your newly created and registered EKI, ERI, and EMI images, use the `euca-describe-images` command. You can also view image status using the Dashboard's **Images** page.

Once you have added your image to Eucalyptus, see [Using Instances](#) section for information about launching and using instances based on the image.

Image Tasks

This section explains the tasks that you can perform on images in Eucalyptus.

You can perform the following image-related tasks listed in the following sections:

- [Add an Image](#)
- [Modify an Image](#)
- [Create an Image](#)
- [Associate a Kernel and Ramdisk](#)
- [Bundle an Image for Amazon EC2](#)
- [Bundle a Windows Instance](#)

Add an Image

An image is the basis for instances that you spin up for your computing needs. If you have access to images that meet your needs, you can skip this section. This section explains how to add an image to your Eucalyptus cloud that is customized for your needs.

There are a few ways you can add an image to Eucalyptus:

- **Add an image based on an existing image.** If you have an image already, read the rest of this section. Eucalyptus has stock images available to help you get started right away. You can find links to these images in the Dashboard's start page. You can also get images from the [Eucalyptus Machine Images](#) page.

- Add an image that you modify from an existing image. If you have an image that you want to modify, see [Modify an Image](#).
- Add an image that you create. If you want to create a new image, see [Create an Image](#).

Each of these three ways has a different first step, but the way you get image to Eucalyptus is the same. To enable an image as an executable entity, you do the following:

1. Bundle a root disk image and kernel/ramdisk pair
2. Upload the bundled image to Walrus bucket storage
3. Register the data with Eucalyptus



Important: Note that while all users can bundle, upload and register images, only the administrator has the required permissions to upload and register kernels and ramdisks.

Once you have an image that meets your needs, perform the tasks listed in this section to add the image to your cloud.

Add a Kernel

When you add a kernel to Walrus, you bundle the kernel file, upload the file to a bucket in Walrus that you name, and then register the kernel with Eucalyptus.

To add a kernel to Walrus:

Use the following three commands:

```
euca-bundle-image -i <kernel_file> --kernel true
euca-upload-bundle -b <kernel_bucket> -m /tmp/<kernel_file>.manifest.xml
euca-register <kernel_bucket>/<kernel_file>.manifest.xml
```

For example:

```
euca-bundle-image -i euca-fedora-10-x86_64/xen-kernel/
vmlinuz-2.6.27.21-0.1-xen --kernel true
...
Generating manifest /tmp/vmlinuz-2.6.27.21-0.1-xen.manifest.xml

euca-upload-bundle -b example_kernel_bucket -m /tmp/
vmlinuz-2.6.27.21-0.1-xen.manifest.xml
...

euca-register example_kernel_bucket/vmlinuz-2.6.27.21-0.1-
xen.manifest.xml
IMAGE eki-XXXXXXX
```

Where the returned value eki-XXXXXXX is the unique ID of the registered kernel image.

Add a Ramdisk

When you add a ramdisk to Walrus, you bundle the ramdisk file, upload the file to a bucket in Walrus that you name, and then register the ramdisk with Eucalyptus.

To add a ramdisk to Walrus:

Use the following three commands:

```
euca-bundle-image -i <ramdisk_file> --ramdisk true
euca-upload-bundle -b <ramdisk_bucket> -m /tmp/<ramdisk_file>.manifest.xml
euca-register <ramdisk_bucket>/<ramdisk_file>.manifest.xml
```

For example:

```
euca-bundle-image -i euca-fedora-10-x86_64/xen-kernel/
initrd-2.6.27.21-0.1-xen --ramdisk true
...
Generating manifest /tmp/initrd-2.6.27.21-0.1-xen.manifest.xml

euca-upload-bundle -b example_rd_bucket -m /tmp/
initrd-2.6.27.21-0.1-xen.manifest.xml
...
Uploaded image as example_rd_bucket/initrd-2.6.27.21-0.1-
xen.manifest.xml

euca-register example_rd_bucket/initrd-2.6.27.21-0.1-
xen.manifest.xml
IMAGE eri-XXXXXXXX
```

Where the returned value `eri-XXXXXXXX` is the unique ID of the registered ramdisk image.

Add a Root Filesystem

When you add a root filesystem to Walrus, you bundle the root filesystem file, upload the file to a bucket in Walrus that you name, and then register the root filesystem with Eucalyptus. The bundle operation can include a registered ramdisk (ERI ID) and a registered kernel (EKI ID). The resulting image would will associate the three images.

You can also bundle the root file system independently and associate the ramdisk and kernel with the resulting EMI at run time. For more information, see [Associate a Kernel and Ramdisk](#).

To add a root filesystem to Walrus:

Use the following three commands:

```
euca-bundle-image -i <root_filesystem_file>
euca-upload-bundle -b <root_filesystem_file_bucket> -m /tmp/
<root_filesystem_file>.manifest.xml
euca-register <root_filesystem_file_bucket>/
<root_filesystem_file>.manifest.xml
```

For example:

```
euca-bundle-image -i euca-fedora-10-x86_64/fedora.10.x86-64.img
--ramdisk eri-722B3CBA --kernel eki-5B3D3859
...
Generating manifest /tmp/fedora.10.x86-64.img.manifest.xml

euca-upload-bundle -b example_rf_bucket -m /tmp/
fedora.10.x86-64.img.manifest.xml
...
Generating manifest /tmp/fedora.10.x86-64.img.manifest.xml

euca-register example_rf_bucket/
fedora.10.x86-64.img.manifest.xml
IMAGE emi-XXXXXXXX
```

Where the returned value `emi-XXXXXXXX` is the unique ID of the registered machine image.

Modify an Image

You might find that existing images in your cloud don't meet your needs.

To modify an existing image:

1. Create a mount point for your image.

```
mkdir temp-mnt
```

2. Associate a loop block device to the image.

```
losetup /dev/loop5 <image_name>
```

where loop5 is a free device

3. Mount the image.

```
mount /dev/loop5 temp-mnt
```

4. Make procfs, dev and sysfs available in your chroot environment.

```
mkdir -p temp-mnt/proc
mkdir -p temp-mnt/sys
mkdir -p temp-mnt/dev
mount -o bind /proc temp-mnt/proc
mount -o bind /sys temp-mnt/sys
mount -o bind /dev temp-mnt/dev
```

5. You now have the image under temp-mnt and you can copy over what you want into it. If you want to install packages into it, you have few options:

- chroot temp-mnt and use apt-get, yum, or zypper to install what you want.
- Instruct the package manager program to use a different root (for example both dpkg and rpm uses the --root option)

6. Unmount the drive.

```
umount /dev/loop5
losetup -d /dev/loop5
```

You now have an image with your modifications. You are ready to [add the image](#) to Eucalyptus.

Create an Image

You can create your own image if you to deploy an instance that is customized for your specific use. For example, you can create an EMI to deploy a LAMP server, MySQL database server, or a Postgres database server. All you need is some familiarity with the creation process.

About Linux Images

A Linux image is a root partition of a Linux installation. We use the convention followed by EC2 images:

- The first partition is the root partition (where the EMI is attached)
- The second is the ephemeral partition (for additional storage)
- The third is the swap partition

Any image you create must conform to these conditions (for example, `/etc/fstab` must follow this convention). Typically, the disk provided by Eucalyptus (as well as EC2) is the first SCSI disk (`sda`). However, you can customize it to be, for example, the first IDE disk (`hda`) or the first virtual disk on a paravirtual machine (`xvda`).

About Windows Images

Eucalyptus supports a licensed version of Windows images. Before you bundle a Windows image, you must have a valid version of Windows OS installed on your hypervisor and the Eucalyptus Windows Integration Service installed in the created Windows VM. The following sections detail how to perform these tasks.

Eucalyptus is compatible with images created from licensed versions of Windows Server 2003 R2, Windows Server 2008, Windows Server 2008 R2, and Windows 7. Windows OS is sensitive to physical and virtual hardware changes made after installation. We recommend that you create any Windows image on the specific hypervisor (Xen, KVM, or VMware) where you plan to run instances of that Windows image. Eucalyptus does not support running a Windows image across different hypervisors.

Before you begin, you will need the following room for a blank disk file:

- 8GB minimum, for Windows Server 2003 R2 (32-bit and 64-bit)
- 15GB minimum for Windows Server 2008 (32-bit and 64-bit), Windows Server 2008 R2 (64-bit), and Windows 7 (32-bit and 64-bit)



Tip: Eucalyptus implements ephemeral disks by which a running Windows instances are allocated an additional disk space based on the instance type. For example, if you assign an instance type 15 GB of disk and the registered Windows EMI is 8 GB, 7 GB of disk spaces are ephemeral disks accessible under D:\. If your Windows application uses scratch space most of time, we recommend keeping the Windows EMI small because it takes longer to launch bigger EMIs.

The windows image is a root file system that has no kernel and ramdisk associated with it. After creating your own Windows images, you can bundle, upload, and register it with Eucalyptus.

Hypervisor and OS Information

The follow sections detail how to create an image based on the hypervisor you use. To create an image, choose which hypervisor you will use, which OS you want the image to be, and then follow the directions for that choice.

Hypervisor	OS	Directions
Xen	Linux	See Create Linux Image (Xen)
Xen	Windows	See Create Windows Image (Xen)
KVM	Linux	See Create Linux Image (KVM)
KVM	Windows	See Create Windows Image (KVM)
VMware	Linux	See Create Linux Image (VMware)
VMware	Windows	See Create Windows Image (VMware)

Create Linux Image (Xen)



Important: Before you begin, make sure that Xen installed and properly functioning for the account you are going to use.

To create a root filesystem to be used with Eucalyptus using Xen and an ISO image of the guest OS you want to install:

1. Download the ISO image of the distro (guest OS) you want to install.
2. Create a virtual disk with the desired size. For example, to create a 4 GiB disk:

```
dd if=/dev/zero of=<image_name> bs=1M count=4096
```

3. Create a Xen configuration file (for example, `xen.cfg`) to boot off the ISO image. The file should look like the following:

```
name = "bootFromISO"
kernel = "/usr/lib/xen/boot/hvmloader"
memory = 1024
builder = "hvm"
device_model = "/usr/lib64/xen/bin/qemu-dm"
boot = "d"
disk = [ 'file:PATH_TO_ISO,hdc:cdrom,r',
'file:PATH_TO/<image_name>,sda,w' ]
vif = [ 'mac=00:01:01:00:00:03, bridge=xenbr0' ]
dhcp="on"
vnc = 1
vncdisplay = 7
pae = 1
```



Important: You must modify `device_model` and `kernel` since they depend on where the distribution puts such files. Look into the Xen package file list. Also be sure that all the options are correct for your needs.

4. Start the domU.

```
xen create xen.cfg
```

5. Connect with a VNC viewer. For example:

```
xvncviewer localhost:7
```

6. Install the distro as you would normally would. Do not use `lvm` or `md` and install everything in one partition. Remember the partition in which the distro is installed.



Important: Most distributions will remember the network interfaces by the MAC address, so interfaces will always be the same (that is, `eth0` and `eth1` will not be inverted accidentally). In Eucalyptus you have to get the distribution to forget the NIC, so that every instance will see the interface as the first interface. To do so you may want to refer to the specific distribution. For example Debian and Ubuntu use `udev` so check `/etc/udev/rules.d/*net*`. CentOS might have something in `/etc/network-scripts/ifcfg-eth0`. Failure to do so may leave your instance without network.

7. Stop the domU.
8. Find out the starting block and the block size of the root file system.

```
parted <image_name>
```

9. Change the units to blocks.

```
(parted) u
Unit? [compact]? b
```

10. Print the current partition table using `p`.

Note the start for the partition you are interested in and the size. For this example we'll say that the start is 32256 and the size is 1024000.

11. Extract the file system.

```
dd if=<image_name> of=rootfs.img bs=1 skip=32256 count=1024000
```

You now have a root filesystem. Make sure that it is compatible with the kernel/initrd you are using in your cloud environment. In particular you may want to be sure you have the modules of the kernel you are going to use. Then you are ready to [add the image](#) to Eucalyptus.

Create Windows Image (Xen)

We recommend that you perform this task on an NC (or the host running the same Linux distributions and hypervisors as the NC). If you are creating the Windows image on a machine currently running the NC, terminate all running instances and stop the NC. To stop the NC, enter:

```
service eucalyptus-nc stop
```

Included with the installation are template files that closely match those that Eucalyptus will generate at VM instantiation time. The file for KVM (`libvirt-xen-windows-example.xml`) is located at `/usr/share/eucalyptus/doc/`. We recommend that you review this file to acquaint yourself with its contents, noting required files, bridges, and resources. For more information about configuring the `libvirt.xml` file, go to the [Domain XML format](#) page in the libvirt documentation.

Install Base Windows OS

The first task for creating a Windows image is installing a base Windows operating system (OS). To install a base Windows OS:

1. Log in to the stopped NC server or a host that runs the same hypervisor as the NCs.
2. Create a blank disk file. Specify your Windows VM image name using the parameter `of`.

```
dd if=/dev/zero of=windows.<image_name>.img bs=1M count=1 seek=16999
```



Important: Your image name must start with the word, windows (all lower-case).

3. Create a floppy and secondary blank disk to be attached to the image later, in order to test paravirtualization drivers

```
dd if=/dev/zero of=floppy.img \
bs=1k count=1474
dd if=/dev/zero of=secondary.img \
bs=1M count=1 seek=1000
```

4. Make available the `.iso` file that contains your Windows installation CD/DVD, then copy the `.iso` file to your working directory.
5. Copy the `libvirt-xen-windows-example.xml` file to your working directory and rename it to `libvirt-xen-windows.xml`.

```
cp libvirt-xen-windows-example.xml /usr/share/eucalyptus/doc/libvirt-xen-
windows.xml
```

6. Open the new `libvirt-xen-windows.xml` file and provide fully qualified paths to the VM image file and iso. Make sure that the name of the bridge is the same as the one used by the hypervisor on which you are creating the Windows image.

Your file should look similar to this example:

```
<domain type='xen'>
  <name>eucalyptus-windows</name>
  <os>
    <type>hvm</type>
    <loader>/usr/lib/xen/boot/hvmloder</loader>
    <boot dev='cdrom' />
  </os>
  <features>
    <acpi/>
    <apic/>
    <pae/>
  </features>
</domain>
```



```

</features>
<clock offset='localtime' />
<on_poweroff>destroy</on_poweroff>
<on_reboot>restart</on_reboot>
<on_crash>destroy</on_crash>
  <memory>524288</memory>
  <vcpu>1</vcpu>
  <devices>
    <emulator>/usr/lib64/xen/bin/qemu-dm</emulator>
    <disk type='file'>
      <source file='/root/windows_2003.img' />
      <target dev='hda' bus='ide' />
    </disk>
    <!--<disk type='file' device='disk'>
      <driver name='tap' type='aio' />
      <source file='fully_qualified_path_to_secondary_disk' />
      <target dev='xvda' bus='xen' />
    </disk>
    <disk type='file' device='floppy'>
      <source file='fully_qualified_path_to_floppy_disk' />
      <target dev='fda' />
    </disk> -->
    <disk type='file' device='cdrom'>
      <source file='/root/
en_win_srv_2003_r2_enterprise_with_sp2_cd1_x13-05460.iso' />
      <target dev='hdc' />
      <readonly />
    </disk>
    <interface type='bridge'>
      <source bridge='xenbr0' />
      <script path='/etc/xen/scripts/vif-bridge' />
    </interface>
    <serial type='pty'>
      <source path='/dev/pts/3' />
      <target port='0' />
    </serial>
    <input type='tablet' bus='usb' />
    <input type='mouse' bus='ps2' />
    <graphics type='vnc' port='-1' autoport='yes' keymap='en-us'
listen='0.0.0.0' />
  </devices>
</domain>

```

7. Start the VM.

```
virsh create libvirt-xen-windows.xml
```

8. Connect to the virtual console using the VNC client of your choice. On the NC, check the display number that has been allocated by looking at the process table (`ps axw | grep vnc`). For example, if the display number is 0, then connect to the NC using the VNC client:

```
vinagre <machine-hosting-vm>:0
```



Tip: You can also connect to your Windows image using the Virtual Manager and Virtual Viewer utilities with Centos/Xen systems.

9. Follow the standard Windows installation procedure until the VM has completed installing Windows.



Tip: On some hosts, the VNC's display number will change when a VM restarts. Use `ps` to find the current number.

10. Run `virsh list` to display the domain name.

11. Shut down the Windows VM you have just created. The easiest way to shutdown your VM is use the `virsh destroy` command, as shown:

```
virsh destroy <domain_name>
```

You are now ready to [Install Eucalyptus Windows Integration](#).

Install Eucalyptus Windows Integration

To install Eucalyptus Windows Integration:

1. Use `yum` or `apt-get` to download the Eucalyptus tool.

```
yum install eucalyptus-windows-tools
apt-get install eucalyptus-windows-tools
```

2. Copy `/usr/share/eucalyptus/eucawin.iso` to the working directory of your NC.

3. Open the `libvirt-xen-windows-example.xml` file and make the following edits:

- Comment out the lines of XML code directing Xen to boot the Windows image from the CDROM.
- Change the text so that `eucawin.iso` replaces the Windows `.iso` image and is mounted as `cdrom`.
- Enter the path to the secondary disk file you created in the [previous task](#).
- Uncomment the lines that direct attachment of a floppy disk and secondary disk.

Your finished file should look similar to the following example:

```
<domain type='xen'>
  <name>eucalyptus-windows</name>
  <os>
    <type>hvm</type>
    <loader>/usr/lib/xen/boot/hvmlloader</loader>
    <!-- <boot dev='cdrom' /> -->
  </os>
  <features>
    <acpi/>
    <apic/>
    <pae/>
  </features>
  <clock offset='localtime' />
  <on_poweroff>destroy</on_poweroff>
  <on_reboot>restart</on_reboot>
  <on_crash>destroy</on_crash>
  <memory>524288</memory>
  <vcpu>1</vcpu>
  <devices>
    <emulator>/usr/lib64/xen/bin/qemu-dm</emulator>
    <disk type='file'>
      <source file='/root/windows_2003.img' />
      <target dev='hda' bus='ide' />
    </disk>
    <disk type='file' device='disk'>
      <driver name='tap' type='aio' />
      <source file='/root/secondary.img' />
      <target dev='xvda' bus='xen' />
    </disk>
    <disk type='file' device='floppy'>
      <source file='/root/floppy.img' />
      <target dev='fda' />
    </disk>
    <disk type='file' device='cdrom'>
      <source file='eucawin.iso' />
      <target dev='hdc' />
    </disk>
  </devices>
</domain>
```

```

        <readonly/>
    </disk>
    <interface type='bridge'>
        <source bridge='xenbr0' />
        <script path='/etc/xen/scripts/vif-bridge' />
    </interface>
    <serial type='pty'>
        <source path='/dev/pts/3' />
        <target port='0' />
    </serial>
    <input type='tablet' bus='usb' />
    <input type='mouse' bus='ps2' />
    <graphics type='vnc' port='-1' autoport='yes' keymap='en-us'
listen='0.0.0.0' />
    </devices>
</domain>

```

4. There is an issue with Xen paravirtualization drivers on some Windows versions. 64-bit Windows Server 2008, Windows Server 2008R2, and Windows 7 have driver signing requirements that prevent unsigned drivers from being installed. The supplied drivers in the package are not signed. To resolve this issue:

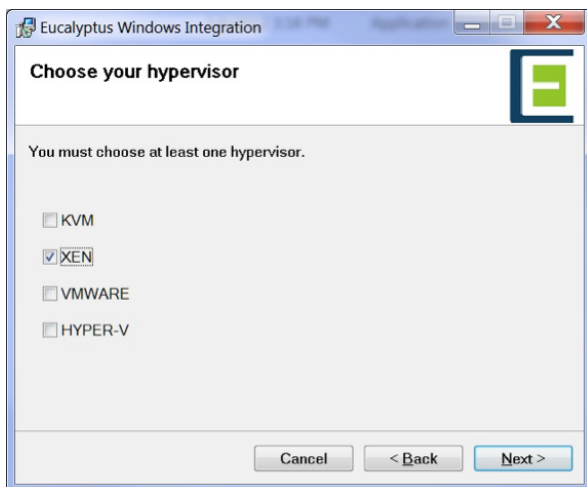
- a) Log in to the Windows VM.
- b) Launch `cmd.exe` as administrator, and execute the following commands:

```

bcdedit.exe -set loadoptions DISABLE_INTEGRITY_CHECKS
bcdedit.exe -set TESTSIGNING ON

```

- c) Reboot the VM.
5. Log in to Windows and find the Eucalyptus installation files in the CDROM drive.
 - For Windows Server 2003 R2, run `setup.exe`.
 - For all others, run `EucalyptusWindowsIntegration.msi`. (`setup.exe` will automatically install .NET framework 2.0, which is not bundled in Server 2003 R2).
6. In the **Choose your hypervisor** step, click **XEN** and then click **Next**.



During the installation of drivers, warning messages might appear asking to confirm the driver installation. Click **Continue** in all cases.

7. At the end of Eucalyptus installation, the Xen driver setup will automatically start. When prompted to choose setup type, choose **Complete**. During the installation of drivers, warning messages might appear asking to confirm the driver installation. Click **Continue** in all cases.
8. Reboot the Windows VM
9. Open the Windows device manager and check that the following drivers are found for each device:
 - Floppy disk drive

- Disk drives: XEN PV SCSI Disk Device
- SCSI and RAID controllers: Xen Block Device Driver
- Network adapters: Xen Net Device Driver

If the correct drivers are not found, question marks display on the devices.

- Right-click on the devices in question and run update drivers. The new hardware wizard displays.
- When the new hardware wizard asks if Windows update is to be connected, click **No, not this time**.
- Choose **Install software automatically (recommended)**.
- If a confirmation popup message displays, click **Continue**.

You are now ready to [Configure Active Directory](#).

Configure Active Directory

The Eucalyptus Integration service lets an enterprise with existing Active Directory domains attach Windows instances to the domains and control access to these instances using the existing AD user database. Users can log into the instance either using their domain credentials or the Administrator's password generated with the `euca-get-password` command.

Because AD technology is tightly integrated with domain name service (DNS), the default name server contacted by the instance must be able to resolve the AD address as a proper domain controller. You can do this for all networking modes except System, by configuring the following line in the CC's `eucalyptus.conf` file:

```
VNET_DNS=<domain_controller_IP_address>
```

If there is no such pre-existing DNS set-up or your networking mode is System, you might need to change the VM's network interface so that the preferred DNS server points to the domain controller.

To set up Active Directory:

1. Click **Windows Programs > Eucalyptus > EucaSetup.exe**.



The **Eucalyptus Windows Integration** popup displays.

2. Click the **Active Directory** tab in the Eucalyptus Windows Integration window and enter the following information:

The screenshot shows the 'Eucalyptus Windows Integration' dialog box with the 'RemoteDesktop' tab selected. The 'General' and 'ActiveDirectory' tabs are also visible. The 'AD Address' field contains 'dc.eucalyptus.com'. The 'Admin Username' field contains 'Eucalyptus'. The 'Admin Password' and 'Confirm' fields both contain '*****'. The 'Organizational Unit (optional)' field contains 'OU=Eucalyptus,DC=eucalyptus,DC=com'. The 'Status' field shows 'not a member of a domain' in red text. At the bottom are 'Apply', 'Close', and 'Clear' buttons.

- Enter the address of the existing Active Directory domain controller in the **AD Address** field.
- Enter the administrator username in the **Admin Username** field. We recommend using a generic user account that has permission to join a computer to a domain or a specific organizational unit.
- Enter and confirm the password in the **Admin Password** field. Note that the Admin username and password are required to join an instance to an Active Directory. When launched in Eucalyptus, these properties will be deleted as soon as the instance joins (or fails to join) the domain.
- Optionally, enter an organizational unit in the **Organizational Unit** field. This specifies a container that the instances launched from this image will be attach to.



Tip: If the values entered in this section are incorrect, the launched instances will fail to join the domain. We recommend that you verify the information by manually joining a computer to a domain using the same information that you entered in this step. You may first log in the launched instance using administrator password (`euca-get-password`) and manually join the domain for verification.

3. Click **Apply**.

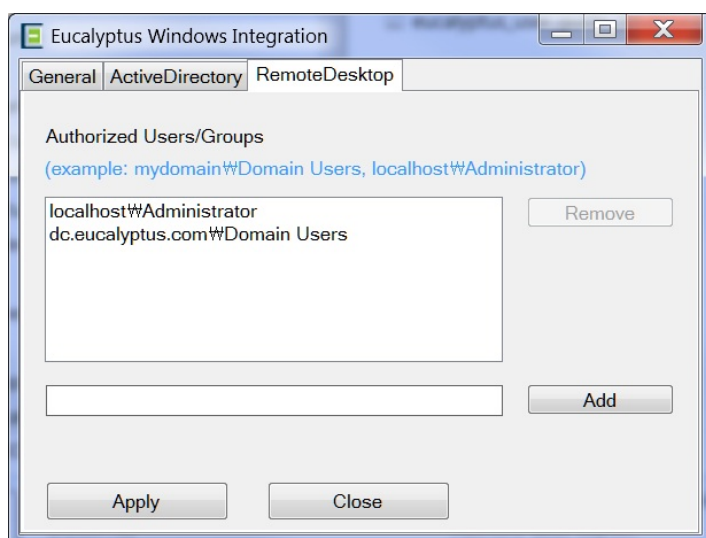
You are now ready to [Configure Remote Desktop](#)

Configure Remote Desktop

Domain users or groups require remote desktop permission to log into an instance. By default, only the local administrator has the remote desktop permission. The Eucalyptus Integration Service provides a way to grant remote desktop permission to additional domain users or groups.

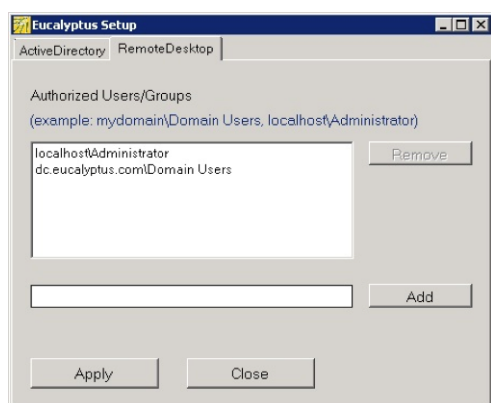
To configure remote desktop permission:

1. Open the **Eucalyptus Windows Integration** popup (**Windows Programs** > **Eucalyptus** > **EucaSetup.exe**).
2. Click the **RemoteDesktop** tab



The names of authorized domain users and groups display in the **Authorized User/Groups** field. By default, only the local administrator is listed as authorized.

3. In the text field below the list, enter a user and group account name in the format [DOMAIN]\[USER or GROUP]. If you add a new local user or local group, prepend the account name localhost\ instead of the domain name.



4. Click **Add**.
5. Repeat for all user/groups that you want to add.
6. Click **Apply**.

When the instance launches, the members of the groups you added can log in to the instance through remote desktop.

You are now ready to [Run Sysprep](#).

Run Sysprep

Sysprep is a Microsoft tool for deploying multiple Windows operating systems in an enterprise. Running Sysprep removes system-specific information such as security ID (SID) from the Windows OS before you clone an image. Sysprep then re-initializes the OS after the image is cloned and started on multiple computers. Use Sysprep to prepare images when you use Microsoft Key Management Service to activate license keys. Also, use Sysprep when your Windows systems are attached to Active Directory.

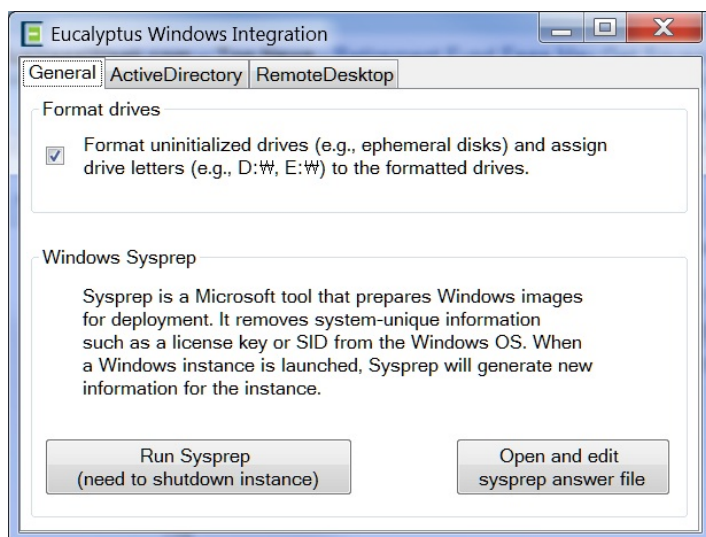
In Eucalyptus, you can run Sysprep before you bundle images with the `euca-bundle-image` or `euca-bundle-instance` commands.



Note: The Eucalyptus Integration Service supports Sysprep for Windows Server 2008, Windows Server 2008 R2, and Windows 7.

To configure and run Sysprep:

1. Open the **Eucalyptus Windows Integration** popup (**Windows Programs > Eucalyptus > EucaSetup.exe**).



2. Click the **General** tab.
3. Ensure that **Format uninitialized drives** is checked
4. If you want to edit the Sysprep answer file, click the **Open and change** button in the General tab. Otherwise skip this step.
5. Click the **Run Sysprep** button.
Sysprep starts.
6. After Sysprep is complete, close the application and shutdown the Windows VM using the Windows Programs menu.

You are now ready to [Add Image to Eucalyptus](#).

Add Image to Eucalyptus

To enable an image as an executable entity, you must bundle and upload the Windows disk image to Walrus, and then register the uploaded image with Eucalyptus.

Run the following command to bundle, upload, and register your Windows disk image:

```
euca-bundle-image -i <vm_image_file>
euca-upload-bundle -b <image_bucket> -m /tmp/<vm_image_file>.manifest.xml
euca-register <image_bucket>/<vm_image_file>.manifest.xml
```

Your Windows image is now ready to run as an instance.

After you register the image, Walrus decrypts the image bundle. This process might take a few minutes for a large Windows image to be decrypted. For example, a 10G image requires that you wait about 10 minutes before you launch the instance.

Create Linux Image (KVM)



Important: Before you begin, make sure that KVM installed and properly functioning for the account you are going to use.

To create a root filesystem to be used with Eucalyptus using KVM and an ISO image of the guest OS you want to install:

1. Download the ISO image of the distro (guest OS) you want to install.

2. Create a virtual disk with the desired size. For example, to create a 4 GiB disk:

```
dd if=/dev/zero of=<image_name> bs=1M count=4096
```

3. Start KVM using the CDROM as the boot device and new image as the disk.

```
kvm -cdrom iso_image -drive if=scsi,file=<image_name>,boot=off
```



Note: You need to have a SCSI as a bus because Eucalyptus expects the disk to be in `/dev/sda`.

4. Install the distro as you would normally would. Do not use `lvm` or `md` and install everything in one partition. Remember the partition in which the distro is installed.



Important: Most distributions will remember the network interfaces by the MAC address, so interfaces will always be the same (that is, `eth0` and `eth1` will not be inverted accidentally). In Eucalyptus you have to get the distribution to forget the NIC, so that every instance will see the interface as the first interface. To do so you may want to refer to the specific distribution. For example Debian and Ubuntu use `udev` so check `/etc/udev/rules.d/*net*`. CentOS might have something in `/etc/network-scripts/ifcfg-eth0`. Failure to do so may leave your instance without network.

5. Find out the starting block and the block size of the root file system.

```
parted <image_name>
```

6. Change the units to blocks.

```
(parted) u
Unit? [compact]? b
```

7. Print the current partition table using `p`.

Note the start for the partition you are interested in and the size. For this example we'll say that the start is 32256 and the size is 1024000.

8. Extract the file system.

```
dd if=<image_name> of=rootfs.img bs=1 skip=32256 count=1024000
```



Tip: To speed up the process, you can increase the block size. Divide the size and the skip count by the new block size (512 in this example) and you get `dd if=<image_name> of=rootfs.img bs=512 skip=63 count=2000`.

You now have a root filesystem. Make sure that it is compatible with the kernel/initrd you are using in your cloud environment. In particular you may want to be sure you have the modules of the kernel you are going to use. Then you are ready to [add the image](#) to Eucalyptus.

Create Windows Image (KVM)

We recommend that you perform this task on an NC (or the host running the same Linux distributions and hypervisors as the NC). If you are creating the Windows image on a machine currently running the NC, terminate all running instances and stop the NC. To stop the NC, enter:

```
service eucalyptus-nc stop
```

Included with the installation are template files that closely match those that Eucalyptus will generate at VM instantiation time. The file for KVM (`libvirt-xen-windows-example.xml`) is located at `/usr/share/eucalyptus/doc/`. We recommend that you review this file to acquaint yourself with its contents, noting required files, bridges, and resources. For more information about configuring the `libvirt.xml` file, go to the [Domain XML format](#) page in the libvirt documentation.

Install Base Windows OS

The first task for creating a Windows image is installing a base Windows operating system (OS). To install a base Windows OS:

1. Log in to the stopped NC server or a host that runs the same hypervisor as the NCs.
2. Create a blank disk file. Specify your Windows VM image name using the parameter `of`.

```
dd if=/dev/zero of=windows.<image_name>.img bs=1M count=1 seek=16999
```



Important: Your image name must start with the word, windows (all lower-case).

3. Create a floppy and secondary blank disk to be attached to the image later, in order to test paravirtualization drivers

```
dd if=/dev/zero of=floppy.img \
bs=1k count=1474
dd if=/dev/zero of=secondary.img \
bs=1M count=1 seek=1000
```

4. Make available the .iso file that contains your Windows installation CD/DVD, then copy the .iso file to your working directory (`/usr/share/eucalyptus/doc/`).
5. Copy the `libvirt-kvm-windows-example.xml` file to your working directory and rename it to `libvirt-kvm-windows.xml`.

```
cp libvirt-kvm-windows-example.xml /usr/share/eucalyptus/doc/libvirt-kvm-
windows.xml
```

6. Open the new `libvirt-kvm-windows.xml` file and provide fully qualified paths to the VM image file and iso. Make sure that the name of the bridge is the same as the one used by the hypervisor on which you are creating the Windows image.

Your file should look similar to this example:

```
<domain type='kvm'>
  <name>eucalyptus-windows</name>
  <os>
    <type>hvm</type>
    <boot dev='cdrom' />
  </os>
  <features>
    <acpi />
  </features>
  <memory>524288</memory>
  <vcpu>1</vcpu>
  <devices>
    <emulator>/usr/bin/kvm</emulator>
    <disk type='file'>
      <source file='/root/windows_2003.img' />
      <target dev='hda' />
    </disk>
    <!-- <disk type='file' device='disk'>
      <source file='fully_qualified_path_to_secondary_disk' />
      <target dev='vda' bus='virtio' />
    </disk>
    <disk type='file' device='floppy'>
      <source file='fully_qualified_path_to_floppy_disk' />
      <target dev='fda' />
    </disk> -->
  </devices>
</domain>
```

```

        <disk type='file' device='cdrom'>
          <source file='/root/
en_win_srv_2003_r2_enterprise_with_sp2_cd1_x13-05460.iso' />
          <target dev='hdc' />
          <readonly />
        </disk>
        <interface type='bridge'>
          <source bridge='br0' />
          <model type='rtl8139' />
        </interface>
        <!--<interface type='bridge'>
          <source bridge='br0' />
          <model type='virtio' />
        </interface> -->
        <graphics type='vnc' port='-1' autoport='yes' listen='0.0.0.0' />
      </devices>
    </domain>

```

7. Start the VM.

```
virsh create libvirt-kvm-windows.xml
```

8. Connect to the virtual console using the VNC client of your choice. On the NC, check the display number that has been allocated by looking at the process table (`ps axw | grep vnc`). For example, if the display number is 0, then connect to the NC using the VNC client:

```
vinagre <machine-hosting-vm>:0
```

9. Follow the standard Windows installation procedure until the VM has completed installing Windows.



Tip: On some hosts, the VNC's display number will change when an image restarts. Use `ps` to find the current number.

10. Run `virsh list` to display the domain name.

11. Shut down the Windows VM you have just created. The easiest way to shutdown your VM is to use the `virsh destroy` command, as shown:

```
virsh destroy <domain_name>
```

You are now ready to [Install Eucalyptus Windows Integration](#).

Install Eucalyptus Windows Integration

To install Eucalyptus Windows Integration:

1. Use `yum` or `apt-get` to download the Eucalyptus tool.

```

yum install eucalyptus-windows-tools
apt-get install eucalyptus-windows-tools

```

2. Copy `/usr/share/eucalyptus/eucawin.iso` to the working directory of your NC.

3. Open the `libvirt-kvm-windows-example.xml` file and make the following edits:

- Comment out the lines of XML code directing KVM to boot the Windows image from the CDROM.
- Change the text so that `eucawin.iso` replaces the Windows `.iso` image and is mounted as `cdrom`.
- Enter the path to the secondary disk file you created in the [previous task](#).
- Uncomment the lines that direct attachment of a floppy disk, secondary disk, and secondary network interface.

Your finished file should look similar to the following example:

```
<domain type="kvm">
```

```

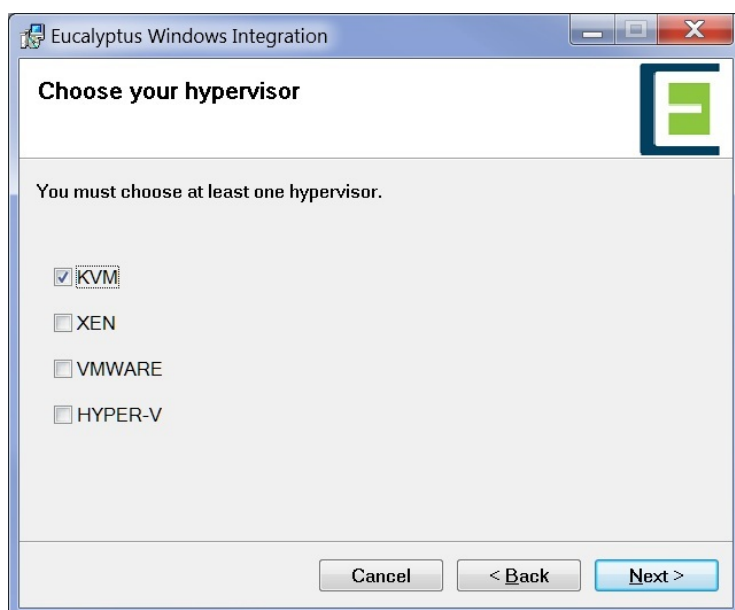
<name>eucalyptus-windows</name>
<os>
  <type>hvm</type>
  <!-- <boot dev='cdrom' /> -->
</os>
<features>
  <acpi/>
</features>
<memory>524288</memory>
<vcpu>1</vcpu>
<devices>
  <emulator>/usr/bin/kvm</emulator>
  <disk type='file'>
    <source file='/root/windows_2003.img' />
    <target dev='hda' />
  </disk>
  <disk type='file' device='disk'>
    <source file='/root/secondary.img' />
    <target dev='vda' bus='virtio' />
  </disk>
  <disk type='file' device='floppy'>
    <source file='/root/floppy.img' />
    <target dev='fda' />
  </disk>
  <disk type='file' device='cdrom'>
    <source file='eucawin.iso' />
    <target dev='hdc' />
    <readonly/>
  </disk>
  <interface type='bridge'>
    <source bridge='br0' />
    <model type='rtl8139' />
  </interface>
  <interface type='bridge'>
    <source bridge='br0' />
    <model type='virtio' />
  </interface>
  <graphics type='vnc' port='-1' autoport='yes' listen='0.0.0.0' />
</devices>
</domain>

```

4. Log in to Windows and find the Eucalyptus installation files in the CDROM drive.

- For Windows Server 2003 R2, run `setup.exe`.
- For all others, run `EucalyptusWindowsIntegration.msi`. (`setup.exe` will automatically install .NET framework 2.0, which is not bundled in Server 2003 R2).

5. In the **Choose your hypervisor** step, click **KVM** and then click **Next**.



During the installation of KVM VirtIO drivers, warning messages might appear asking to confirm the driver installation. Click **Continue** in all cases.

6. At the **Welcome to the Found New Hardware Wizard**, click **Next**.
7. Choose **Install Software Automatically...** and click **Next**.
8. Reboot the Windows VM
9. Open the Windows device manager and check that the following drivers are found for each device:
 - Floppy disk drive
 - Disk drivers: Red Hat VirtIO SCSI Disk Device
 - SCSI and RAID controllers: Red Hat VirtIO SCSI controller
 - Network adapters: Red Hat VirtIO Ethernet Adapter

If the correct drivers are not found, question marks display on the devices.

- Right-click on the devices in question and run update drivers. The new hardware wizard displays.
- When the new hardware wizard asks if Windows update is to be connected, click **No, not this time**.
- Choose **Install software automatically (recommended)**.
- If a confirmation popup message displays, click **Continue**.

You are now ready to [Configure Active Directory](#).

Configure Active Directory

The Eucalyptus Integration service lets an enterprise with existing Active Directory domains attach Windows instances to the domains and control access to these instances using the existing AD user database. Users can log into the instance either using their domain credentials or the Administrator's password generated with the `euca-get-password` command.

Because AD technology is tightly integrated with domain name service (DNS), the default name server contacted by the instance must be able to resolve the AD address as a proper domain controller. You can do this for all networking modes except System, by configuring the following line the CC's `eucalyptus.conf` file:

```
VNET_DNS=<domain_controller_IP_address>
```

If there is no such pre-existing DNS set-up or your networking mode is System, you might need to change the VM's network interface so that the preferred DNS server points to the domain controller.

To set up Active Directory:

1. Click **Windows Programs > Eucalyptus > Eucalyptus Setup**.



The **Eucalyptus Windows Integration** popup displays.

2. Click the **Active Directory** tab in the Eucalyptus Windows Integration window and enter the following information:

- Enter the address of the existing Active Directory domain controller in the **AD Address** field.
- Enter the administrator username in the **Admin Username** field. We recommend using a generic user account that has permission to join a computer to a domain or a specific organizational unit.
- Enter and confirm the password in the **Admin Password** field. Note that the Admin username and password are required to join an instance to an Active Directory. When launched in Eucalyptus, these properties will be deleted as soon as the instance joins (or fails to join) the domain.
- Optionally, enter an organizational unit in the **Organizational Unit** field. This specifies a container that the instances launched from this image will be attach to.



Tip: If the values entered in this section are incorrect, the launched instances will fail to join the domain. We recommend that you verify the information by manually joining a computer to a domain using the same information that you entered in this step. You may first log in the launched instance using administrator password (`euca-get-password`) and manually join the domain for verification.

3. Click **Apply**.

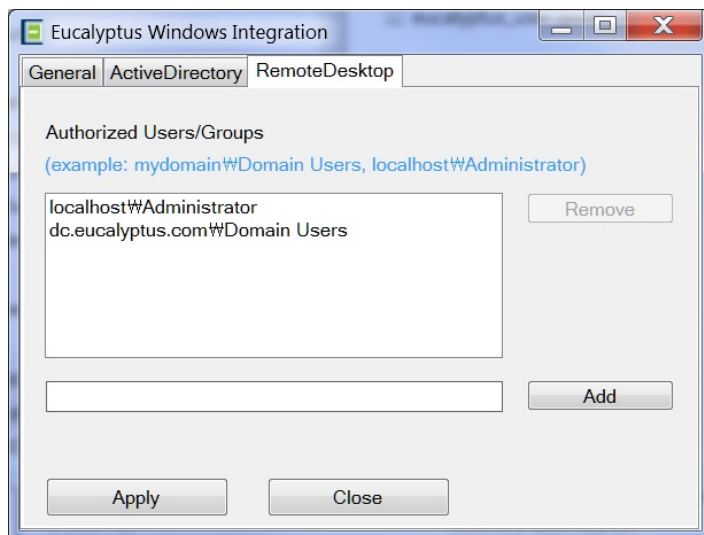
You are now ready to [Configure Remote Desktop](#).

Configure Remote Desktop

Domain users or groups require remote desktop permission to log into an instance. By default, only the local administrator has the remote desktop permission. The Eucalyptus Integration Service provides a way to grant remote desktop permission to additional domain users or groups.

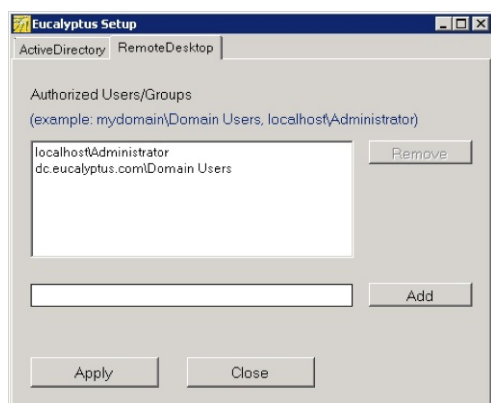
To configure remote desktop permission:

1. Open the **Eucalyptus Windows Integration** popup (**Windows Programs > Eucalyptus > Eucalyptus Setup**).
2. Click the **RemoteDesktop** tab



The names of authorized domain users and groups display in the **Authorized User/Groups** field. By default, only the local administrator is listed as authorized.

3. In the text field below the list, enter a user and group account name in the format [DOMAIN]\[USER or GROUP]. If you add a new local user or local group, prepend the account name localhost\ instead of the domain name.



4. Click **Add**.
5. Repeat for all user/groups that you want to add.
6. Click **Apply**.

When the instance launches, the members of the groups you added can log in to the instance through remote desktop.

You are now ready to [Run Sysprep](#).

Run Sysprep

Sysprep is a Microsoft tool for deploying multiple Windows operating systems in an enterprise. Running Sysprep removes system-specific information such as security ID (SID) from the Windows OS before you clone an image.

Sysprep then re-initializes the OS after the image is cloned and started on multiple computers. Use Sysprep to prepare images when you use Microsoft Key Management Service to activate license keys. Also, use Sysprep when your Windows systems are attached to Active Directory.

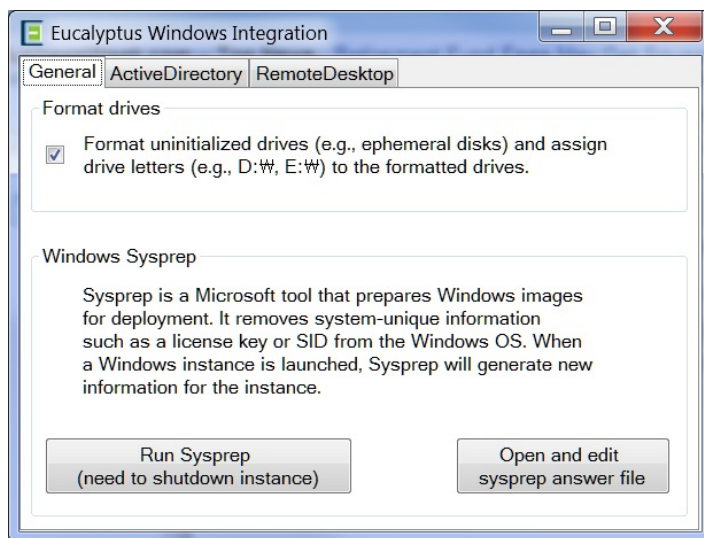
In Eucalyptus, you can run Sysprep before you bundle images with the `euca-bundle-image` or `euca-bundle-instance` commands.



Note: The Eucalyptus Integration Service supports Sysprep for Windows Server 2008, Windows Server 2008 R2, and Windows 7.

To configure and run Sysprep:

1. Open the **Eucalyptus Windows Integration** popup (**Windows Programs** > **Eucalyptus** > **Eucalyptus Setup**).



2. Click the **General** tab.
3. Ensure that **Format uninitialized drives** is checked
4. If you want to edit the Sysprep answer file, click the **Open and change** button in the General tab. Otherwise skip this step.
5. Click the **Run Sysprep** button.
Sysprep starts.
6. After Sysprep is complete, close the application and shutdown the Windows VM using the Windows Programs menu.

You are now ready to [Add Image to Eucalyptus](#).

Add Image to Eucalyptus

To enable an image as an executable entity, you must bundle and upload the Windows disk image to Walrus, and then register the uploaded image with Eucalyptus.

Run the following command to bundle, upload, and register your Windows disk image:

```
euca-bundle-image -i <vm_image_file>
euca-upload-bundle -b <image_bucket> -m /tmp/<vm_image_file>.manifest.xml
euca-register <image_bucket>/<vm_image_file>.manifest.xml
```

Your Windows image is now ready to run as an instance.

After you register the image, Walrus decrypts the image bundle. This process might take a few minutes for a large Windows image to be decrypted. For example, a 10G image requires that you wait about 10 minutes before you launch the instance.

Create Linux Image (VMware)

To create a root filesystem to be used with Eucalyptus using VMware and an ISO image of the guest OS you want to install:

1. Run an instance of the image.
2. Attach a volume to the instance.

```
euca-attach-volume -i <instance_id> -d <local_device_name> <volume_id>
```

3. Make a file system on the volume and mount it to /mnt.

```
mkfs -t ext3 /dev/vdb
mount /dev/vdb /mnt
```

4. SCP your Eucalyptus credential zip file to the instance.

```
scp -i yourkey.private euca2-xxxxxxx-x509.zip root@xxx.xxx.xxx.xxx:/mnt
```

5. SSH into the instance.

```
ssh -i yourkey.private root@xxx.xxx.xxx.xxx
```

6. Install the tools you need.

```
yum install rsync unzip
```

7. Change to the /mnt device.

```
cd /mnt
```

8. Unzip your Eucalyptus credentials zip to /mnt/euca.

```
unzip euca2-admin-x509.zip -d euca/
```

9. Install any other software you want on the instance so it will be part of your new image.
10. Remove the udev rules.

```
rm /etc/udev/rules.d/70*
```

You now have a root filesystem. Make sure that it is compatible with the kernel/initrd you are using in your cloud environment. In particular you may want to be sure you have the modules of the kernel you are going to use. Then you are ready to [add the image](#) to Eucalyptus.

Create Windows Image (VMware)

To create an image from a Windows OS in VMware you will need one network interface and one disk. This disk must be the first IDE disk. Follow the tasks listed in this section.

Install the Base Windows OS

To install the base Windows operating system (OS):

1. Create a new VM using the VMware vSphere Client following the standard VMware procedures.



Note: Make sure that the Windows VM uses the LSI Logic Parallel driver as the SCSI controller. For some Windows versions, it is not the default SCSI controller in the VM setting.

2. Install the Windows OS following standard VMware procedures.
3. Install VMware Tools.

You are now ready to [Install Eucalyptus Windows Integration](#).

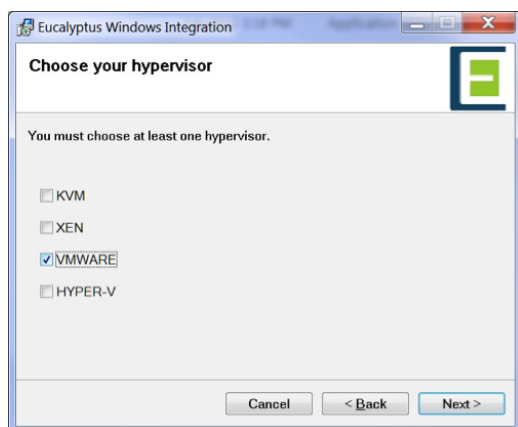
Install Eucalyptus Windows Integration

To install the Eucalyptus Windows Integration Service:

1. Use yum or apt-get to download the Eucalyptus tool.

```
yum install eucalyptus-windows-tools
apt-get install eucalyptus-windows-tools
```

2. Copy `$EUCALYPTUS/usr/share/eucalyptus/eucawin.iso` to the Windows host that runs vSphere client. You might need SFTP client on the Windows host.
3. Use the VMware vSphere client to upload the ISO file to the VSphere datastore.
4. Attach the `eucawin.iso` to the Windows VM.
5. Log in to Windows and find the Eucalyptus installation files in the CDROM drive.
 - For Windows Server 2003 R2, run `setup.exe`. This automatically installs the .NET framework 2.0, which is not bundled in Server 2003 R2.
 - For all other versions, run `EucalyptusWindowsIntegration.msi`.
6. In the **Choose your hypervisor** step, click **VMWARE**.



7. Click **Next** and continue until the end of installation.

You are now ready to [Configure Active Directory](#).

Configure Active Directory

The Eucalyptus Integration service lets an enterprise with existing Active Directory domains attach Windows instances to the domains and control access to these instances using the existing AD user database. Users can log into the instance either using their domain credentials or the Administrator's password generated with the `euca-get-password` command.

Because AD technology is tightly integrated with domain name service (DNS), the default name server contacted by the instance must be able to resolve the AD address as a proper domain controller. You can do this for all networking modes except System, by configuring the following line the CC's `eucalyptus.conf` file:

```
VNET_DNS=<domain_controller_IP_address>
```

If there is no such pre-existing DNS set-up or your networking mode is System, you might need to change the VM's network interface so that the preferred DNS server points to the domain controller.

To set up Active Directory:

1. Click **Windows Programs > Eucalyptus > Eucalyptus Setup**.
The **Eucalyptus Windows Integration** popup displays.
2. Click the **Active Directory** tab in the Eucalyptus Windows Integration window and enter the following information:

The screenshot shows the 'Eucalyptus Windows Integration' dialog box with the 'ActiveDirectory' tab selected. The fields are filled with the following values:

- AD Address:** dc.eucalyptus.com
- Admin Username:** Eucalyptus
- Admin Password:** ***** (masked)
- Confirm:** ***** (masked)
- Organizational Unit (optional):** OU=Eucalyptus,DC=eucalyptus,DC=com
- Status:** not a member of a domain

At the bottom, there are three buttons: 'Apply', 'Close', and 'Clear'.

- Enter the address of the existing Active Directory domain controller in the **AD Address** field.
- Enter the administrator username in the **Admin Username** field. We recommend using a generic user account that has permission to join a computer to a domain or a specific organizational unit.
- Enter and confirm the password in the **Admin Password** field. Note that the Admin username and password are required to join an instance to an Active Directory. When launched in Eucalyptus, these properties will be deleted as soon as the instance joins (or fails to join) the domain.
- Optionally, enter an organizational unit in the **Organizational Unit** field. This specifies a container that the instances launched from this image will be attach to.



Tip: If the values entered in this section are incorrect, the launched instances will fail to join the domain. We recommend that you verify the information by manually joining a computer to a domain using the same information that you entered in this step. You may first log in the launched instance using administrator password (`euca-get-password`) and manually join the domain for verification.

3. Click **Apply**.

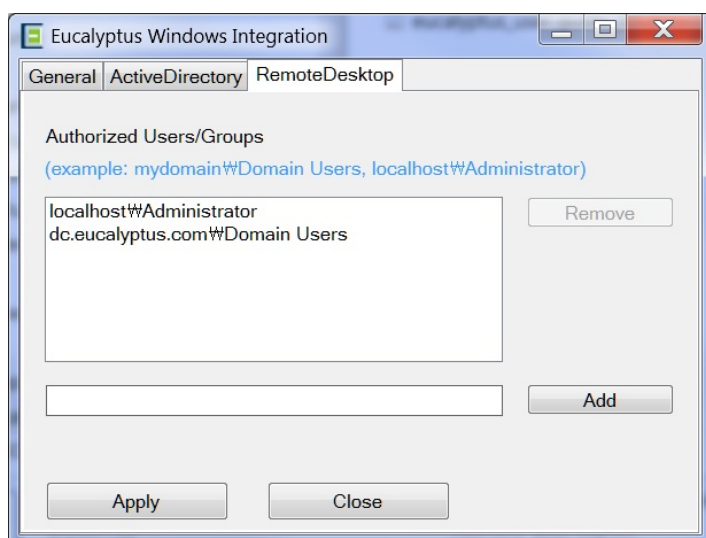
You are now ready to [Configure Remote Desktop](#).

Configure Remote Desktop

Domain users or groups require remote desktop permission to log into an instance. By default, only the local administrator has the remote desktop permission. The Eucalyptus Integration Service provides a way to grant remote desktop permission to additional domain users or groups.

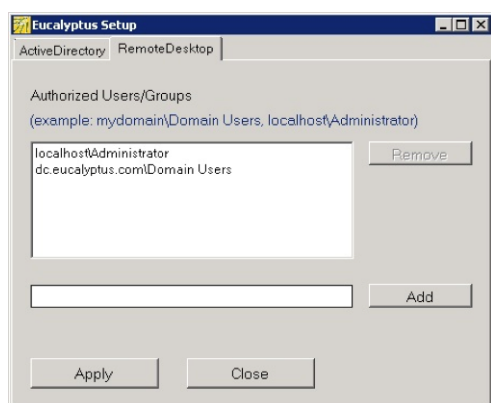
To configure remote desktop permission:

1. Open the **Eucalyptus Windows Integration** popup (**Windows Programs** > **Eucalyptus** > **Eucalyptus Setup**).
2. Click the **RemoteDesktop** tab



The names of authorized domain users and groups display in the **Authorized User/Groups** field. By default, only the local administrator is listed as authorized.

3. In the text field below the list, enter a user and group account name in the format [DOMAIN]\[USER or GROUP]. If you add a new local user or local group, prepend the account name localhost\ instead of the domain name.



4. Click **Add**.
5. Repeat for all user/groups that you want to add.
6. Click **Apply**.

When the instance launches, the members of the groups you added can log in to the instance through remote desktop.

You are now ready to [Run Sysprep](#).

Run Sysprep

Sysprep is a Microsoft tool for deploying multiple Windows operating systems in an enterprise. Running Sysprep removes system-specific information such as security ID (SID) from the Windows OS before you clone an image. Sysprep then re-initializes the OS after the image is cloned and started on multiple computers. Use Sysprep to prepare images when you use Microsoft Key Management Service to activate license keys. Also, use Sysprep when your Windows systems are attached to Active Directory.

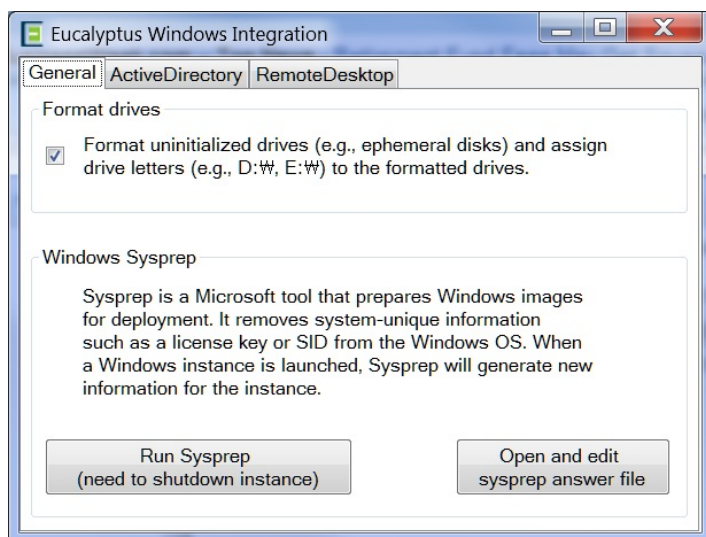
In Eucalyptus, you can run Sysprep before you bundle images with the `euca-bundle-image` or `euca-bundle-instance` commands.



Note: The Eucalyptus Integration Service supports Sysprep for Windows Server 2008, Windows Server 2008 R2, and Windows 7.

To configure and run Sysprep:

1. Open the **Eucalyptus Windows Integration** popup (**Windows Programs > Eucalyptus > Eucalyptus Setup**).



2. Click the **General** tab.
3. Ensure that **Format uninitialized drives** is checked
4. If you want to edit the Sysprep answer file, click the **Open and change** button in the General tab. Otherwise skip this step.
5. Click the **Run Sysprep** button.
Sysprep starts.
6. After Sysprep is complete, close the application and shutdown the Windows VM using the Windows Programs menu.

You are now ready to [Convert the VMDK to an Image](#).

Convert the VMDK to an Image

After your Windows VM image is ready on ESX/ESX-i/VCenter, use the `euca_imager` command to download the image and convert it to the format Eucalyptus uses. The `euca_imager` command converts a remote VMDK located on a datastore into a disk image on the local disk on the CC/VMware Broker machine.

1. Log into the CC/VMware Broker machine and, if necessary, set the appropriate environment variables.

```
# export VDDK_HOME=/opt/packages/vddk/
# export EUCALYPTUS=/
# export LD_LIBRARY_PATH=/opt/packages/vddk/lib/vmware-vix-disklib/lib64/
# PATH=/usr/lib/eucalyptus/:$PATH
# cd ~/.euca (or path to .eucarc)
# source .eucarc
```

2. Enter the `euca_imager` command with appropriate parameters to extract and convert the VMDK to a raw disk image. Parameters are described in the following table:

Parameter	Description
debug	Enables <code>euca_imager</code> debugging output
in	https URL for the vCenter Server or ESX/ESXi machine
out	User-defined name of disk
login	Name used to access the VMware machine
password	Password used to access the VMware machine

Parameter	Description
vsphere-vmrk	Go to vSphere client, find the virtual machine that you have installed (make sure it is in powered off state), and select the Hard disk 1 resource. There will be a Disk File field for the VMDK, and a similar field on the Option tab for the VMX file.
vsphere-vmx	Same description as vsphere-vmrk parameter. This is optional if you're contacting an ESX node directly, but may be required when contacting vCenter.
vsphere-datacenter	Specified name of the datacenter. This is optional if you're contacting an ESX node directly, but may be required when contacting vCenter.

The following example shows euca_imager with the appropriate parameters as described above converting a VMDK to a raw disk image. Note that the prefix of the out parameter must begin with the word windows.

```
/usr/lib/eucalyptus/euca_imager debug=yes convert in=
https://192.168.7.198 out=windows.disk login=Administrator
password=password
vsphere-vmrk="[datastore1]WindowsVM/WindowsVM.vmrk" vsphere-
datacenter=DataCenter1
vsphere-vmx="[datastore1]WindowsVM/WindowsVM.vmx"
```

You are now ready to [Add Image to Eucalyptus](#).

Add Image to Eucalyptus

To enable an image as an executable entity, you must bundle and upload the Windows disk image to Walrus, and then register the uploaded image with Eucalyptus.

Run the following command to bundle, upload, and register your Windows disk image:

```
euca-bundle-image -i <vm_image_file>
euca-upload-bundle -b <image_bucket> -m /tmp/<vm_image_file>.manifest.xml
euca-register <image_bucket>/<vm_image_file>.manifest.xml
```

Your Windows image is now ready to run as an instance.

After you register the image, Walrus decrypts the image bundle. This process might take a few minutes for a large Windows image to be decrypted. For example, a 10G image requires that you wait about 10 minutes before you launch the instance.

Associate a Kernel and Ramdisk

There are three ways that you can associate a kernel and ramdisk with an image.



Tip: This only applies to Linux images. Windows images do not use kernels or ramdisks.

- You can associate a specific kernel and ramdisk identifier with an image at the euca-bundle-image step.

```
euca-bundle-image -i <vm_image_file> --kernel <eki-XXXXXXX> --ramdisk
<eri-XXXXXXX>
```

- You can specify kernel and ramdisk at instance run time as an option to `euca-run-instances` command.

```
euca-run-instances --kernel <eki-XXXXXXXX> --ramdisk <eri-XXXXXXXX> <emi-XXXXXXXX>
```

- An administrator can set default registered kernel and ramdisk identifiers that will be used if a kernel and ramdisk are unspecified by either of the other options.

Bundle an Image for Amazon EC2

To bundle an image for use with Amazon EC2:

1. Locate the Amazon ec2 cert file that is provided as part of the EC2 AMI tools. This file is generally located in `$EC2_AMITOOL_HOME/etc/ec2/amitools/cert-ec2.pem`.
2. Enter the following command:

```
euca-bundle-image -i <image_name> \
-c <cert_filename> -k <private_key_filename> \
--ec2cert <path_to_cert_file>
```

For example:

```
euca-bundle-image -i euca-centos-5.3 \
x86_64/centos.5-3.x86-64.img -u 123456789111 -c cert- \
abc.pem -k pk-abc.pem --ec2cert \
$EC2_AMITOOL_HOME/etc/ec2/amitools/cert-ec2.pem
```

Bundle a Windows Instance

The `euca-bundle-instance` command lets you bundle a new Windows image from a running Windows instance directly to Walrus storage. Using `euca-bundle-instance` is an efficient way to generate modified Windows VMs. You can spin up a Windows VM instance from an existing Windows VM image, modify it as needed, then save the modified image to Walrus storage, where it is immediately available for registering and running with the modifications already in place.

To bundle a Windows instance:

Enter the following command:

```
euca-bundle-instance -b <bucket_name> -p
<prefix_starting_with_windows> -o $EC2_ACCESS_KEY -w
$EC2_SECRET_KEY <instance_ID>
```



Tip: You can query the progress of a bundling task using the `euca-describe-bundle-task` command. You can cancel the active bundle task using `euca-cancel-bundle-task` command.

The following example bundles a new Windows image from an existing Windows instance with ID `i-12c4af6a`, to the bucket `mybucket`, with image prefix `windows2003_bundle`:

```
euca-bundle-instance -b mybucket -p windows2003_bundle -o
$EC2_ACCESS_KEY -w $EC2_SECRET_KEY i-12c4af6a
```

Appendix

This appendix for the Eucalyptus Administration Guide provides information about administrative commands and general troubleshooting help.

Eucalyptus Administration Commands

Eucalyptus offers commands for common administration tasks and inquiries. This section provides a reference for these commands.

euca_conf

Short reference description.

Syntax

```
euca_conf
```

Options

Option	Description	Required
--initialize	Begin the one-time initialization of the CLC	No
--heartbeat	Return heartbeat data for the specified host	No
--synckey		No
--no-rsync	Do not use rsync when registering	No
--no-scp	Do not use scp when registering	No
--skip-scp-hostcheck	Skip scp interactive host keycheck	No
--get-credentials	Download credentials to the specified zip file. By default, the admin credentials will be downloaded but this can be adjusted with the --cred-user option. Each time this is called, new X.509 certificates will be created for the specified user.	No
--cred-account	Set get-credentials for the specified account	No
--cred-user	Set get-credentials for the specified user	No
--register-nodes	Add specified nodes to Eucalyptus	No
--deregister-nodes	Remove specified nodes from Eucalyptus	No
--register-arbitrator	Add arbitrator service to Eucalyptus	No
--deregister-arbitrator	Remove arbitrator service from Eucalyptus	No
--register-cloud	Add new Cloud Controller to Eucalyptus	No
--deregister-cloud	Remove Cloud Controller from Eucalyptus	No
--register-cluster	Add a Cluster Controller to Eucalyptus	No
--deregister-cluster	Remove a Cluster Controller from Eucalyptus	No

Option	Description	Required
--register-walrus	Add Walrus to Eucalyptus	No
--deregister-walrus	Remove Walrus from Eucalyptus	No
--register-sc	Add Storage Controller to Eucalyptus	No
--deregister-sc	Remove Storage Controller from Eucalyptus	No
--register-vmwarebroker	Add VMware Broker to Eucalyptus	No
--deregister-vmwarebroker	Remove VMware Broker from Eucalyptus	No
--list-walruses	Return all registered Walruses	No
--list-clouds	Return all registered Cloud Controllers	No
--list-clusters	List all registered Cluster Controllers	No
--list-arbitrators	Return all registered arbitrator services	No
--list-vmwarebrokers	Return all registered VMware Brokers	No
--list-nodes	Return all registered Node Controllers	No
--list-components	return all registered Eucalyptus components	No
--list-services	Return all registered services	No
--list-scs	Return all registered Storage Controllers	No
--no-sync	Used with --register-* to skip syncing keys	No
-d	Point Eucalyptus to the specified directory	No
--cc-port	Set the Cluster Controller to the specified port	No
--sc-port	Set the Storage Controller to the specified port	No
--walrus-port	Set Walrus to the specified port	No
--nc-port	Set the Node Controller to the specified port	No
--instances	Set the instance path	No
--hypervisor	Set which hypervisor to use. Valid values: xen kvm	No
--user	Set the user to use for EUCA_USER	No
--dhcpd	Set the DHCP daemon binary to the specified path	No
--dhcp_user	Set the specified user name to run dhcpd as	No
--bridge	Set the bridge as the specified name	No
--name	Returns the value for the specified name	No
--import-conf	Import variables from a specified eucalyptus.conf file	No
--upgrade-conf	Upgrade eucalyptus.conf from the specified older installation file	No
--setup	Perform initial setup	No
--enable	Enable specified service at next start Valid values: cloud walrus sc vmwarebroker	No

Option	Description	Required
<code>--disable</code>	Disable specified service at next start Valid values: cloud walrus sc vmwarebroker	No
<code>--check</code>	Pre-flight checks Valid values: common vmware	No
<code>-P, --partition</code>	Name of partition. Used with <code>--register-*</code> and <code>--deregister-*</code>	No
<code>-H, --host</code>	Name or IP address of host. Used with <code>--register-*</code>	No
<code>-C, --component</code>	Name of the component. Used with <code>--register-*</code> and <code>--deregister-*</code>	No
<code>--help-register</code>	Display help for register deregister	No

Common Options

Parameter	Description
<code>--region</code> <i>region</i>	Region to direct requests to. Only valid for EC2 endpoints.
<code>-U, --url url</code>	URL of the cloud to connect to. Expects an EC2 endpoint <code>/services/Eucalyptus</code> .
<code>-a, --access-key</code> <i>access_key_id</i>	User's access key ID
<code>-s, --secret-key</code> <i>secret_key</i>	User's secret key
<code>--config</code> <i>configuration_path</i>	Read credentials and cloud settings from the specified config file. Default: <code>\$HOME/.eucarc</code> or <code>/etc/euca2ools/eucarc</code> .
<code>--debug</code>	Prints what the command sends to the server and what it receives from the server. Use when you're trying to debug Euca2ools.
<code>--debugger</code>	Enable interactive debugger on error
<code>-h, --help</code>	Display the manual page for the command.
<code>--version</code>	Display the version of this tool

euca-describe-properties

Lists properties.

Syntax

```
euca-describe-properties
```

Options

None

Common Options

Parameter	Description
<code>--region region</code>	Region to direct requests to. Only valid for EC2 endpoints.
<code>-U, --url url</code>	URL of the cloud to connect to. Expects an EC2 endpoint <code>/services/Eucalyptus</code> .
<code>-a, --access-key access_key_id</code>	User's access key ID
<code>-s, --secret-key secret_key</code>	User's secret key
<code>--config configuration_path</code>	Read credentials and cloud settings from the specified config file. Default: <code>\$HOME/.eucarc</code> or <code>/etc/euca2ools/eucarc</code> .
<code>--debug</code>	Prints what the command sends to the server and what it receives from the server. Use when you're trying to debug Euca2ools.
<code>--debugger</code>	Enable interactive debugger on error
<code>-h, --help</code>	Display the manual page for the command.
<code>--version</code>	Display the version of this tool



euca-describe-services

Returns services.

Syntax

```
euca-describe-services
```

Options

Option	Description	Required
<code>-A, --all</code>	Include all public service information. Reported state information is determined by the view available to the target host, which should be treated as advisory (See documentation for guidance on interpreting this information).	No
<code>--system-internal</code>	Include internal services information  Note: This information is only for the target host.	No
<code>--user-services</code>	Include services that are user-facing and co-located with some other top-level service  Note: This information is only for the target host.	No
<code>-T, --filter-type</code>	Filter services by specified component type	No
<code>-H, --filter-host</code>	Filter services by specified host	No

Option	Description	Required
-F, --filter-state	Filter services by state	No
-P, --filter-partition	Filter services by specified partition	No
-E, --events	Return service event details	No
-events-verbose	Return verbose service event details	No

Common Options

Parameter	Description
--region <i>region</i>	Region to direct requests to. Only valid for EC2 endpoints.
-U, --url <i>url</i>	URL of the cloud to connect to. Expects an EC2 endpoint /services/Eucalyptus.
-a, --access-key <i>access_key_id</i>	User's access key ID
-s, --secret-key <i>secret_key</i>	User's secret key
--config <i>configuration_path</i>	Read credentials and cloud settings from the specified config file. Default: \$HOME/.eucarc or /etc/euca2ools/eucarc.
--debug	Prints what the command sends to the server and what it receives from the server. Use when you're trying to debug Euca2ools.
--debugger	Enable interactive debugger on error
-h, --help	Display the manual page for the command.
--version	Display the version of this tool

Troubleshooting Eucalyptus

Provides common management or configuration problems and solutions for fixing these problems.

To troubleshoot Eucalyptus, the administrator must know the location of the Eucalyptus components, that is, on which machine each component is installed. The administrator must have root access to each machine hosting the components and must know the network configuration connecting the components.

Log Files

Usually when an issue arises in Eucalyptus, you can find a clue or trace or record that suggests the nature of the problem either in the eucalyptus log files or in the system log files. Assuming Eucalyptus is installed in root (/), the eucalyptus logs are located on each machine hosting a component in the following directory: /var/log/eucalyptus/.

Here are the relevant logs for each component:

CLC, Walrus, SC, and VMware Broker

- cloud-debug.log

- `cloud-error.log`
- `cloud-output.log`
- `cloud-exhaust.log`
- `cloud-extreme.log`
- `euca_imager.log`

CC

- `cc.log`
- `axis2c.log`
- `httpd-cc_error_log`

NC

- `cc.log`
- `axis2c.log`
- `httpd-cc_error_log`
- `euca_test_nc.log`

System

You might also find helpful information about the nature of an issue in the system's logs. In particular, search for clues in

- `/var/log/messages`
- `/var/log/libvirt/`
- `/var/log/xen/`

Network Information

When you have to troubleshoot, it's important to understand the elements of the network on your system. Here are some ideas for finding out information about your network:

- It is also important to understand the elements of the network on your system. For example, you might want to list bridges to see which devices are enslaved by the bridge. To do this, use the `brctl` command.
- You might also want to list network devices and evaluate existing configurations. To do this, use these commands: `ip`, `ifconfig`, and `route`.
- If you are running Eucalyptus in Managed networking mode, you can also use `vconfig` to evaluate VLAN configuration.
- You can get further information if you use the `euca-describe` commands with the `verbose` options. For example, `euca-describe-instances verbose` returns all instances running by all users on the system. Other describe commands are:
 - `euca-describe-volumes verbose`
 - `euca-describe-snapshots verbose`
 - `euca-describe-groups verbose`
 - `euca-describe-keypairs verbose`

Walrus and Storage

Walrus decryption failed.

On Ubuntu 10.04 LTS, kernel version 2.6.32-31 includes a bug that prevents Walrus from decrypting images. This can be determined from the following line in `cloud-output.log`

```
javax.crypto.
```

```
BadPaddingException: pad block
corrupted
```

If you are running this kernel:

1. Update to kernel version 2.6.32-33 or higher.
2. De-register the failed image (`euca-deregister`).
3. Re-register the bundle that you uploaded (`euca-register <bucket>/<manifest>`).

Walrus physical disk is not large enough.

1. Stop the CLC.
2. Add a disk.
3. Migrate your data.

Make sure you use LVM with your new disk drive(s).

Access and Identities

Need to verify an existing LIC file.

1. Enter the following command:

```
/usr/sbin/euca-describe-properties
| grep ldap
```

The output from the example above shows the name of the LIC file and status of the synchronization (set to false).

```
PROPERTY
authentication.ldap_integration_configuration
{ 'sync': { 'enable': 'false' } }
```

Windows Images

Properties

A typical size of Windows images is large and Eucalyptus has a set of properties that limit the size of various storage components. The first step in troubleshooting is to make sure that the values are large enough to store your Windows images. You can modify a property using

```
/usr/sbin/euca-modify-property -p <property>=<value>
```

The properties that might affect registering Windows images are:

- `walrus.storage.maxbucketsizeinmb`: max bucket size enforced by Walrus; should be larger than a Windows image
- `walrus.storage.maxcachesizeinmb`: total size of all images that is cached in Walrus; should be larger than the sum of all images (Windows/Linux) in Walrus
- `walrus.storage.maxtotalsnapshotsizingb`: if a Windows image is a type of EBS-backed EMI, this should be large enough to store all EBS backed images
- `{PARTITION}.storage.maxvolumesizingb`: if a Windows image is a type of EBS-backed EMI, this should be large enough to store the image

In addition, during the `euca-run-instances`, the CLC may time-out an instance while a large windows image (images in both Walrus and EBS) is being launched. We recommend that you raise the values of the following properties.

- `cloud.vmstate.instance_timeout`: maximum wait time, in minutes, before the instance becomes running. An instance cannot stay in pending longer than this. Default: 60
- `cloud.vmstate.ebs_volume_creation_timeout`: maximum wait time, in minutes, before a volume backing a boot from EBS image is created. Default: 30
- `cloud.addresses.maxkillorphans`: The public IP assigned to an instance will be expired after the time limit. The exact time-out is $\{\text{maxkillorphans} * 5\}$ seconds (by default it's 50 seconds). If the volume backing an EBS image is not created in time, the public IP will be released from the instance.

Image Preparation

`euca-bundle-image` hangs

The time to bundle an image is proportional to the image size. Because the typical size of Windows image is big, give enough time until bundling is complete. As a rule of thumb, it may take up to 20 min. for bundling a 10 GB Windows image.

`euca-upload-bundle` fails

Make sure `'walrus.storage.maxbucketsizeinmb'` is large enough. If not, ask your administrator.

Instance Launch and Login

Instance stays in pending

Typically, it takes longer to launch Windows images than Linux images as the delay is proportional to the image size. This can be especially long when the image is seeded on NCs the first time (images are cached in NCs and run within few seconds thereafter). As a rule of thumb, 10 GB Windows images may take up to 10 minutes to become 'running' when it is not cached in NCs.

Instance stay in pending and goes to shutdown

An instance may time-out if the Windows image is too big. Review and adjust the relevant properties.

Instance is running, but not accessible using Remote Desktop.

after instances become running, you should wait until Windows is fully booted. If the image is sysprepped, the booting time may take up to 10 min. Also you should make sure the followings are cleared:

- The port 3389 is authorized in the security group
- If the instance is attached to your active directory domain, the domain GPO shouldn't block the RDP port (3389)
- The username should be authorized to log-in using Remote Desktop (refer to User guide: Windows integration service)

Finding the login username and password

Use `Administrator` and the password retrieved by `euca-get-password`. If the instance is attached to a domain, you may use your domain username and password (make sure the username is prepended with domain name, such as `YOUR_DOMAIN\Alice`).

Can't retrieve windows password using `euca-get-password`

Make sure the platform field of your windows EMI is set to 'windows', not 'linux' (use `euca-describe-images`). If not, the most likely reason is that the image name does not begin with 'windows'. You should bundle/upload/register the image with a proper name.

Instance is not attached to an Active Directory domain

- Make sure the parameters set in Windows integration service are correct. One way to verify them is to log in the instance using Administrator password and manually attach the instance to the domain (System Properties -> Computer Name) using the same parameters.
- Make sure VNET_DNS in eucalyptus.conf is set to the domain controller (refer to User Guide: Configure Active Directory).

Disk and Volume

Ephemeral disks are not visible in the Windows

Open Disk Management console (**All Programs->Administrative Tools->Server Manager->Storage**) and find the uninitialized disks. You should create a partition on the disk and format it.

EBS volume is attached, but not visible in the Windows

Open Disk Management console (**All Programs->Administrative Tools->Server Manager->Storage**) and find the uninitialized disks. You should create a partition on the disk and format it. You don't have to repeat it when the volume is reattached later.

EBS volume is detached, but the disk drive (for example, E:\) is still visible in the Windows

For KVM hypervisor, you should perform "remove hardware safely" before detaching the volume.

euca-bundle-instance fails

Make sure the bucket specified with '-b' option doesn't already exist and the property 'walrus.storage_max_bucket_size_inmb' is large enough to store the image.

Instances

Inaccurate IP addresses display in the output of euca-describe-addresses.

This can occur if you add IPs from the wrong subnet into your public IP pool, do a cleanrestart on the CC, swap out the wrong ones for the right ones, and do another clean restart on the CC. To resolve this issue, run the following:

```
/etc/init.d/eucalyptus-cloud stop
/etc/init.d/eucalyptus-cc stop
iptables -F
/etc/init.d/eucalyptus-cc
cleanrestart
/etc/init.d/eucalyptus-cloud start
```

Recovering from a Fail: Walrus

In these examples, we will assume that Walrus WS00 is the primary and WS01 is the secondary Walrus server.

Software Failure Example

In this scenario, WS01 refuses to go to DISABLED state. DRBD complains that it is in split brain mode. drbdadm cstate r0 shows that DRBD is in WFConnection state.

If you are sure that data on WS01 is out of date and can be discarded, execute the following commands to restore HA mode.

1. Shut down the eucalyptus-cloud process on WS01.
2. Ensure that the DRBD connection is down by typing "drbdadm disconnect r0" on any of the two Walrus hosts.
3. On the primary Walrus, WS00, set drbd as the primary by executing "drbdadm primary r0"
4. On the secondary Walrus, WS01, execute the following command:

```
drbdadm -- --discard-my-data connect
```



Warning: This command will discard all data on WS01 and synchronize data from WS00.

5. Monitor the state of DRBD by running:

```
watch -n 2 cat /proc/drbd
```

6. When the data on WS01 is synced, start the eucalyptus-cloud process on WS01.

Hardware Failure Example

In this example, the primary WS00 needs to be taken out of service due to a hardware failure, such as a failed disk.

1. Shut down the eucalyptus-cloud process on WS00 if it is still running.
2. Monitor service status by running `euca-describe-services` on WS01 and ensure that WS01 has taken over as the new primary (state: ENABLED).
3. Shut down the host running WS00.
4. If the host running WS00 is to be replaced entirely or the OS reinstalled:

- On the primary CLC, enter the following to deregister WS00:

```
euca_conf --deregister-walrus --component WS00 partition <name of
partition>
--host <WS00 host>
```

- After Linux has been installed on the new WS00 host and it is ready for use, please reinstall the "eucalyptus-walrus" package.
 - Synchronize the DRBD configuration (/etc/drbd.conf and /etc/eucalyptus/drbd*) from the WS01 host.
 - On WS00, re-configure DRBD by following the Configure DRBD section of the Installation Guide and performing the steps that are relevant to the secondary Walrus server (WS00 is the new secondary Walrus server, in this example).
 - Re-register WS00 with a new host name if necessary. This will synchronize keys.
5. On WS00, execute the following command:

```
drbdadm -- --discard-my-data connect
```



Warning: This command will discard all data on WS00 and synchronize data from WS01.

6. Monitor the state of DRBD by entering:

```
watch -n 2 cat /proc/drbd
```

WS01 should be marked as the primary and WS00 is the new secondary. Wait until data is synchronized.

7. When the data on WS00 is synced from WS01, start the eucalyptus-cloud process on WS00.
8. Monitor service status by running "euca-describe-services" on the primary CLC and ensure that WS00 is DISABLED and WS01 is ENABLED.

At this point, the Walrus service is back in HA mode.

Glossary

cluster

A group of resources that contains a CC, an SC and, optionally, a VMware Broker.

availability zone

An availability zone for AWS denotes a large subset of their cloud environment. Eucalyptus refines this definition to denote a subset of the cloud that shares a local area network. Each availability zone has its own cluster controller and storage controller.

AWS

Amazon Web Services

bucket storage

A storage container that accepts objects via PUT and GET commands.

bundling

A virtual machine image splits the image into multiple image parts to facilitate ease of uploading. It also generates an XML manifest file containing metadata referencing the image, including image parts and kernel, which is used to assemble instances of the image.

Cloud Controller

The Cloud Controller (CLC) is the entry-point into the cloud for administrators, developers, project managers, and end-users. The CLC queries the node managers [SM1] for information about resources, makes high-level scheduling decisions, and makes requests to the Cluster Controllers (CCs). As the interface to the management platform, the CLC is responsible for exposing and managing the underlying virtualized resources (servers, network, and storage). You can access the CLC through Amazon's Elastic Compute Cloud (EC2) and through a web-based Dashboard.

Cluster Controller

The Cluster Controller (CC) generally executes on a machine that has network connectivity to both the machines running the Node Controller (NC) and to the machine running the CLC. CCs gather information about a set of node machines and schedules virtual machine (VM) execution on specific nodes. The CC also manages the virtual machine networks and participates in the enforcement of SLAs[SM3] as directed by the CLC. All nodes associated with a single CC must be in the same broadcast domain (Ethernet).

kernel/ramdisk pair

A ramdisk contains drivers that direct the kernel to launch appropriate system files when instantiating a virtual machine.

Node Controller

The Node Controller (NC) executes on any machine that hosts VM instances. The NC controls VM activities, including the execution, inspection, and termination of VM instances. It also fetches and maintains a local cache of instance images, and it queries and controls the system software (host OS and the hypervisor) in response to queries and control requests from the CC. The NC is also responsible for the management of the virtual network endpoint.

Storage Controller

The Storage Controller (SC) provides functionality similar to the [5] Amazon Elastic Block Storage (EBS) and is capable of interfacing with various storage systems (NFS, iSCSI, SAN devices, etc.). Elastic block storage exports storage volumes that can be attached by a VM and mounted or accessed as a raw block device. EBS volumes persist past VM termination and are commonly used to store persistent data. An EBS volume cannot be shared between VMs and can only be accessed within the same availability zone in which the VM is running. Users can create snapshots from EBS volumes. Snapshots are stored in Walrus and made available across availability zones. Eucalyptus with SAN support lets you use your enterprise-grade SAN devices to host EBS storage within a Eucalyptus cloud.

VMware Broker

VMware Broker (Broker) is an optional Eucalyptus component activated only in versions of Eucalyptus with VMware support. VMware Broker enables Eucalyptus to deploy VMs on VMware infrastructure elements and mediates all interactions between the Cluster Controller (CC) and VMware hypervisors (ESX/ESXi) either directly or through VMware vCenter.

Walrus

Walrus allows users to store persistent data, organized as buckets and objects. You can use Walrus to create, delete, and list buckets, or to put, get, and delete objects, or to set access control policies. Walrus is interface compatible with Amazon's Simple Storage Service (S3), providing a mechanism for storing and accessing virtual machine images and user data. Note that Walrus access is global to the entire Eucalyptus cloud. This means that it can be accessed by end-

users, whether the user is running a client from outside the cloud or from a virtual machine instance running inside the cloud.